

## 1 AGGREGATE OPERATIONS

SQL supports five aggregate operations, which can be applied on any column, say A, of a relation:

- `COUNT ([DISTINCT] A)` : The number of (unique) values in the A column.
- `SUM ([DISTINCT] A)` : The sum of all (unique) values in the A column.
- `AVG ([DISTINCT] A)` : The average of all (unique) values in the A column.
- `MAX (A)` : The maximum value in the A column
- `MIN (A)` : The minimum value in the A column

1: Count the number of different sailor names

```
SELECT COUNT (DISTINCT S.sname)
FROM Sailors S
```

## 2 The GROUP BY and HAVING Clauses

```
SELECT [DISTINCT] select-list
FROM from-list
WHERE qualification
GROUP BY grouping-list
HAVING group-qualification
```

So far we have applied aggregate operations to all rows in a relation. We often want to apply aggregate operations to each of a number of **groups** of a row in a relation.

2: Find the age of the youngest sailor for each rating level

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
GROUP BY S.rating
```

- `GROUP By X` means **put all those with the same value for X in the one group**
- `GROUP By X, Y` means **put all those with the same values for both X and Y in the one group**
- Each column that appears in **select-list** must also appear in the **grouping-list**. This is because each row in the result of the query corresponds to one *group*, which is a collection of rows that agree on values of columns in **grouping-list**.

- If `GROUP BY` is omitted, the entire table is regarded as a single group
- `HAVING` applies a selection criteria to the groups. It can be used to select groups which satisfy a given condition. Think of it as a `WHERE` clause for the returned groups
- A column appearing in the **group-qualification** for `HAVING` must also appear in **grouping-list** for `GROUP BY`
- In general `WHERE` clause is evaluating before `HAVING`. The `WHERE` clause acts as a pre filter where as `HAVING` clause as a post filter.
- The difference between `HAVING` and `WHERE` in SQL is that the `WHERE` cannot be used with aggregates, but `HAVING` can. The `WHERE` clause works on row's data, not on aggregated data.

### 3 EXERCISES

Tables used in this note:

Sailors(sid: integer, sname: string, rating: integer, age: real);

Boats(bid: integer, bname: string, color: string);

Reserves(sid: integer, bid: integer, day: date).

- (1) For each red boat, find the number of reservations for this boat.
- (2) Find the average age of sailors for each rating level that has at least two sailors.
- (3) Find age of the youngest sailor with age  $\geq 18$ , for each rating with at least 2 sailors between 18 and 60.
- (4) Find the average age of sailors who are of voting age (i.e. at least 18 yrs old) for each rating level that has at least two sailors.
- (5) Find the ratings for which the average age of sailors is the minimum over all ratings.

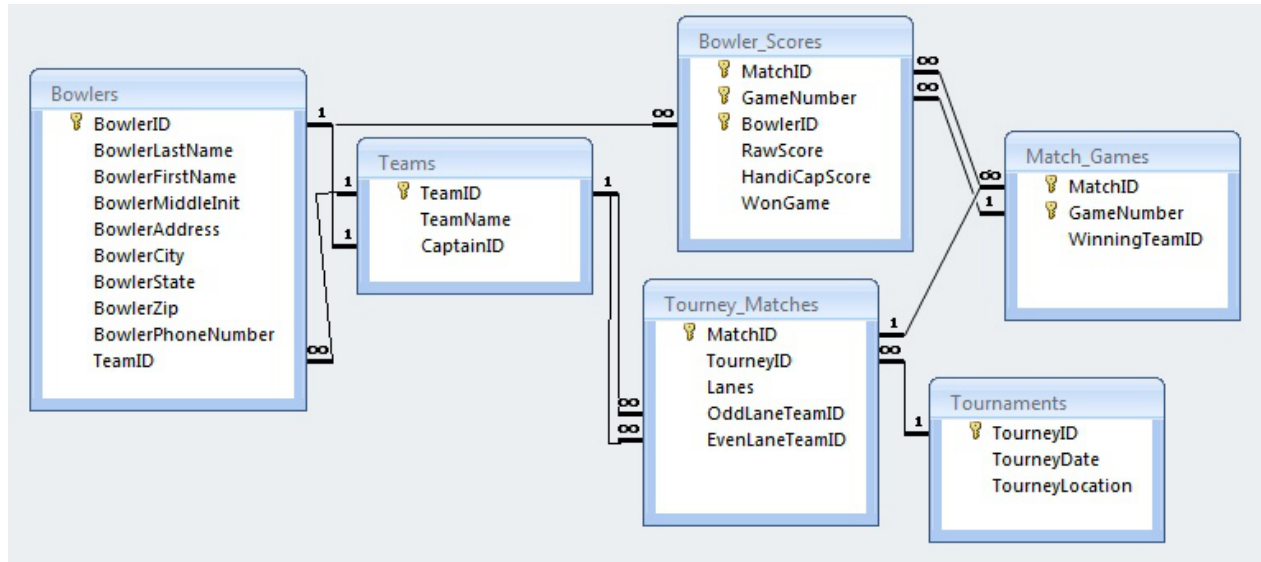
### 4 INNER Joins

```
table1 INNER JOIN table2 ON table1.col1 = table2.col2
```

- Creates a virtual table matching rows to the columns specified after `ON`.
- `INNER` can be omitted.
- If a row is cannot be matched in the other table, then it does not appear in the result virtual table.
- Multiple matches will creat a row for *every pair* of matches.

## 5 EXERCISES

*BowlingLeague.sqlite:*



- (1) Print a schedule of all the team matchups over the whole season (Date, Location, Odd-TeamName, EvenTeamName).
- (2) Print game results for Tournament 1, including match ID, game number, team names, bowler names, and raw score.

## 6 The Memory/Storage Hierarchy

