

# MLDS-413 Introduction to Databases and Information Retrieval

## Lecture 8 Having predicates, INNER Joins

Instructor: Nikos Hardavellas

Slides adapted from Steve Tarzia

# Last Lecture

SELECT description and examples in SQLite

SELECT steps (abbreviated):

1. FROM chooses the table of interest
2. WHERE throws out irrelevant rows
3. GROUP BY identifies rows to combine
4. SELECT tells what values to return (math and aggregation on each group)
5. HAVING throws out irrelevant rows (after aggregation)
6. ORDER BY sorts
7. LIMIT throws out rows based on their position in the results

# Predicates in more detail

- **WHERE & HAVING** filter rows according to conditions called *predicates*
- Any of the following can be combined, like an algebraic expression:
  - Binary operations (used between two things):
    - = == != <> > < >= <= LIKE AND OR REGEXP ←(coming soon!)
    - + - \* / || % << >> & |
    - See [https://www.sqlite.org/lang\\_expr.html](https://www.sqlite.org/lang_expr.html)
  - NOT ...
  - ... IS NULL, ... IS NOT NULL ←(coming soon!)
  - ... BETWEEN ... AND ...
  - ... IN (... , ... , ...)
  - (...)
- Can also use all of the above in the columns we print out, and inside aggregations like **SUM, MIN, MAX, AVG**

# HAVING is like WHERE, but applied after aggregation

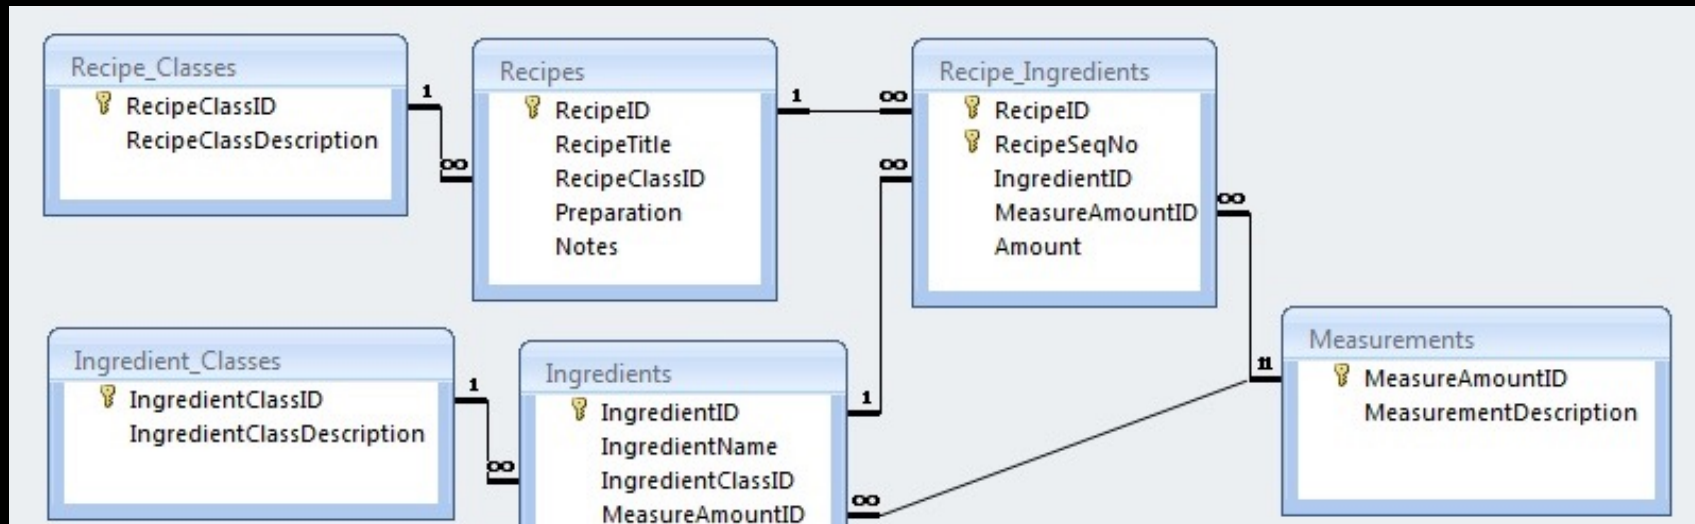
SELECT steps (abbreviated):

1. FROM chooses the table of interest
2. WHERE throws out irrelevant rows
3. GROUP BY identifies rows to combine
4. SELECT tells what values to return (allowing math and aggregation)
  - COUNT, SUM, MAX, MIN, AVG, etc, (i.e., all aggregators) are all evaluated at this step
5. HAVING throws out irrelevant rows (after aggregation)
  - HAVING can now use COUNT, SUM, MAX, MIN, AVG, etc, in conditionals
6. ORDER BY sorts
7. LIMIT throws out rows based on their position in the results

# HAVING example

- Which recipes have 10 or more ingredients (ordered by #ingredients)?

```
SELECT recipeID, count(DISTINCT IngredientID) as NumIngredients
FROM Recipe_Ingredients
GROUP BY recipeID
HAVING NumIngredients >= 10
ORDER BY count(DISTINCT IngredientID);
```



# What if you need to combine data from multiple tables?

1. **FROM** chooses the table of interest
2. **WHERE** throws out irrelevant rows
3. **GROUP BY** identifies rows to combine
4. **SELECT** tells what values to return (allowing math and aggregation)
5. **HAVING** throws out irrelevant rows (after aggregation)
6. **ORDER BY** sorts
7. **LIMIT** throws out rows based on their position in the results

A subquery can draw data from another table, but there is a better way ...

# JOINS create *virtual* tables from several tables

- Normalizing the staff directory left us with three tables
- Eliminated redundant information, but now we have to look in three different tables to answer some questions.

staff			
<i>id</i>	<i>name</i>	<i>room</i>	<i>departmentID</i>
11	Bob	100	1
20	Betsy	100	2
21	Fran	101	1
22	Frank	102	4
35	Sarah	200	5
40	Sam	10	7
54	Pat	102	2

department		
<i>id</i>	<i>name</i>	<i>buildingID</i>
1	Industrial Eng.	1
2	Computer Sci.	2
4	Chemistry	1
5	Physics	4
7	Materials Sci.	5

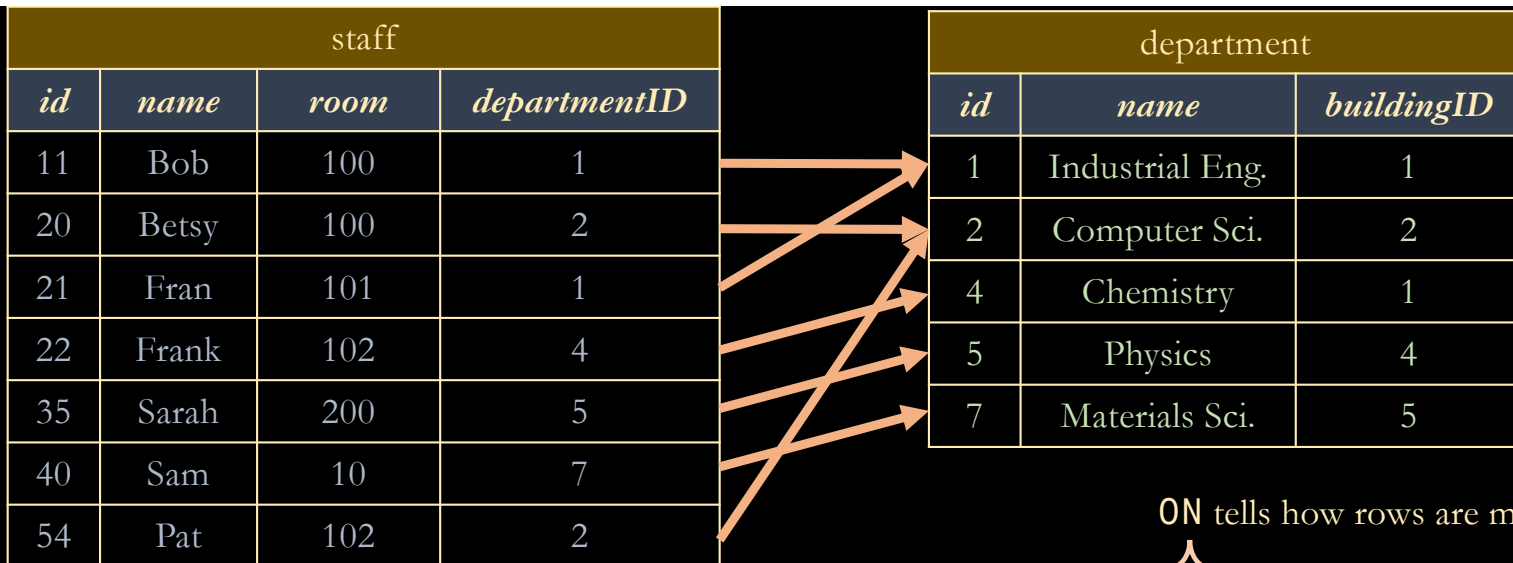
building		
<i>id</i>	<i>name</i>	<i>facilitiesExt</i>
1	Tech	1-1000
2	Ford	1-5003
4	Mudd	1-2005
5	Cook	1-3004
6	Garage	1-6001

# What if we want to print the staff directory?

staff					
<i>id</i>	<i>name</i>	<i>department</i>	<i>building</i>	<i>room</i>	<i>facilitiesExt</i>
11	Bob	Industrial Eng.	Tech	100	1-1000
20	Betsy	Computer Sci.	Ford	100	1-5003
21	Fran	Industrial Eng.	Tech	101	1-1000
22	Frank	Chemistry	Tech	102	1-1000
35	Sarah	Physics	Mudd	200	1-2005
40	Sam	Materials Sci.	Cook	10	1-3004
54	Pat	Computer Sci.	Ford	102	1-5003

We can generate a virtual table like this with INNER JOIN





ON tells how rows are matched

SELECT \* FROM staff JOIN department ON staff.departmentId=department.id

staff. <i>id</i>	staff. <i>name</i>	staff. <i>room</i>	staff. <i>departmentId</i>	department. <i>id</i>	department. <i>name</i>	department. <i>buildingID</i>
11	Bob	100	1	1	Industrial Eng.	1
20	Betsy	100	2	2	Computer Sci.	2
21	Fran	101	1	1	Industrial Eng.	1
22	Frank	102	4	4	Chemistry	1
35	Sarah	200	5	5	Physics	4
40	Sam	10	7	7	Materials Sci.	5
54	Pat	102	2	2	Computer Sci.	2

# How JOIN builds a composite table

```
SELECT * FROM staff JOIN department
ON staff.departmentID=department.id;
```

Start with the first table (staff)

Join with rows from the 2<sup>nd</sup> table (department)  
that match according to the **ON** columns



staff.id	staff.name	staff.room	staff.departmentID	department.id	department.name	department.buildingID
11	Bob	100	1	1	Industrial Eng.	1
20	Betsy	100	2	2	Computer Sci.	2
21	Fran	101	1	1	Industrial Eng.	1
22	Frank	102	4	4	Chemistry	1
35	Sarah	200	5	5	Physics	4
40	Sam	10	7	7	Materials Sci.	5
54	Pat	102	2	2	Computer Sci.	2

# Just print the columns we need

```
SELECT staff.id, staff.name, room,  
       department.name, department.buildingId  
FROM staff JOIN department  
ON staff.departmentId=department.id;
```

staff. <i>id</i>	staff. <i>name</i>	staff. <i>room</i>	department. <i>name</i>	department. <i>buildingID</i>
11	Bob	100	Industrial Eng.	1
20	Betsy	100	Computer Sci.	2
21	Fran	101	Industrial Eng.	1
22	Frank	102	Chemistry	1
35	Sarah	200	Physics	4
40	Sam	10	Materials Sci.	5
54	Pat	102	Computer Sci.	2

# Reorder and rename the columns

```
SELECT staff.id AS staffID,  
       staff.name AS name,  
       department.name AS department,  
       department.buildingId AS buildingId,  
       staff.room AS room  
FROM staff JOIN department  
ON staff.departmentId=department.id;
```

<i>staffId</i>	<i>name</i>	<i>department</i>	<i>buildingID</i>	<i>room</i>
11	Bob	Industrial Eng.	1	100
20	Betsy	Computer Sci.	2	100
21	Fran	Industrial Eng.	1	101
22	Frank	Chemistry	1	102
35	Sarah	Physics	4	200
40	Sam	Materials Sci.	5	10
54	Pat	Computer Sci.	2	102

# Add the third table

```
SELECT staff.id AS staffId, staff.name AS name,  
       department.name AS department,  
       building.name AS building, staff.room AS room,  
       building.faxNumber AS faxNumber  
FROM staff JOIN department ON staff.departmentId = department.id  
       JOIN building ON department.buildingId = building.id;
```

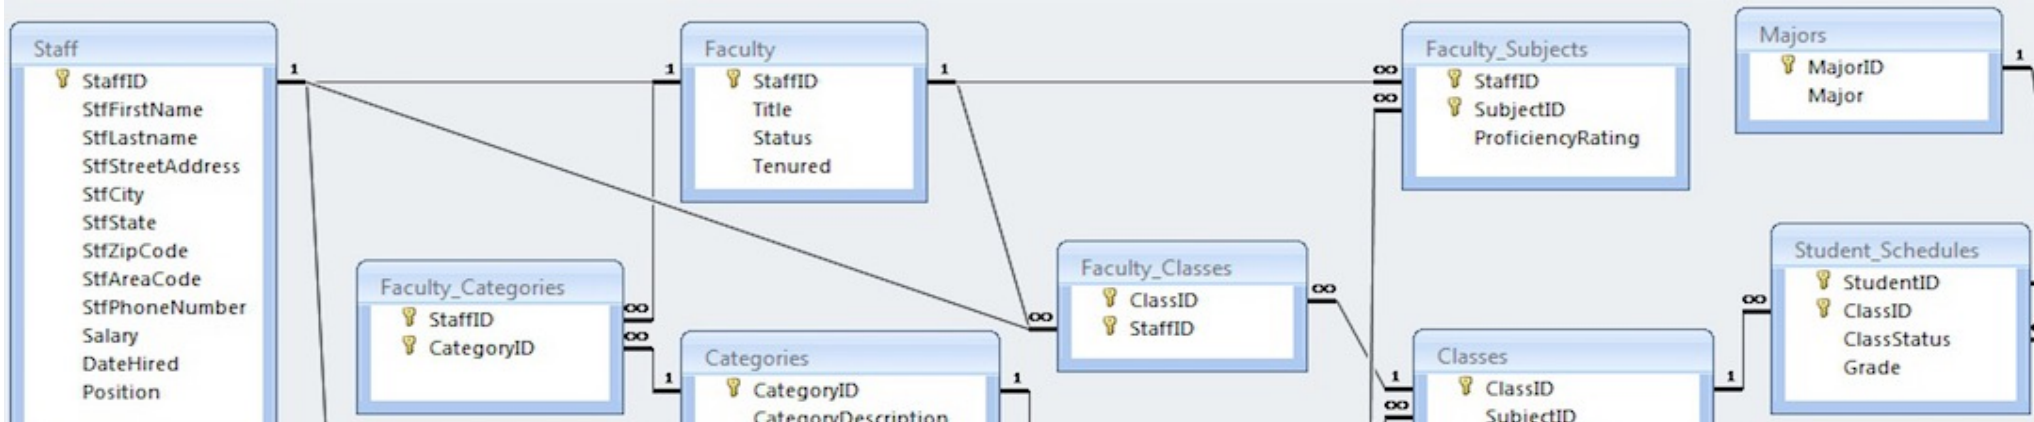
<i>staffId</i>	<i>name</i>	<i>department</i>	<i>building</i>	<i>room</i>	<i>faxNumber</i>
11	Bob	Industrial Eng.	Tech	100	1-1000
20	Betsy	Computer Sci.	Ford	100	1-5003
21	Fran	Industrial Eng.	Tech	101	1-1000
22	Frank	Chemistry	Tech	102	1-1000
35	Sarah	Physics	Mudd	200	1-2005
40	Sam	Materials Sci.	Cook	10	1-3004
54	Pat	Computer Sci.	Ford	102	1-5003

# Who teaches the largest class and what is the average grade?

```
SELECT Student_Schedules.ClassID,  
       AVG(Grade),  
       StfLastname
```

```
FROM Student_Schedules JOIN Faculty_Classes  
     ON Student_Schedules.ClassID=Faculty_Classes.ClassID  
     JOIN Staff ON Faculty_Classes.StaffID = Staff.StaffID  
GROUP BY Student_Schedules.ClassID  
ORDER BY COUNT(*) DESC LIMIT 1;
```

ClassID	AVG(Grade)	StfLastname
2907	78.202	Waldal



# How JOIN builds a composite table

```
SELECT * FROM staff JOIN department
      ON staff.departmentID=department.id
```

What if there  
are multiple  
matches in  
the second  
table?

Start with the first table (staff)

Join with rows from the 2<sup>nd</sup> table (department)  
that match according to the ON columns

staff.id	staff.name	staff.room	staff.departmentID	department.id	department.name	department.buildingID
11	Bob	100	1	1	Industrial Eng.	1
20	Betsy	100	2	2	Computer Sci.	2
21	Fran	101	1	1	Industrial Eng.	1
22	Frank	102	4	4	Chemistry	1
35	Sarah	200	5	5	Physics	4
40	Sam	10	7	7	Materials Sci.	5
54	Pat	102	2	2	Computer Sci.	2

# How JOIN deals with multiple matches

```
SELECT * FROM department JOIN staff  
ON staff.departmentID=department.id
```

department		
<i>id</i>	<i>name</i>	<i>buildingID</i>
1	Industrial Eng.	1
2	Computer Sci.	2
4	Chemistry	1
5	Physics	4
7	Materials Sci.	5

staff			
<i>id</i>	<i>name</i>	<i>room</i>	<i>departmentID</i>
11	Bob	100	1
20	Betsy	100	2
21	Fran	101	1
22	Frank	102	4
35	Sarah	200	5
40	Sam	10	7
54	Pat	102	2



# How JOIN deals with multiple matches

```
SELECT * FROM department JOIN staff
ON staff.departmentId=department.id
```

Create a row for *every pair* of matches


What if there  
are multiple  
matches in  
the second  
table?

department. <i>id</i>	department. <i>name</i>	department. <i>buildingId</i>	staff. <i>id</i>	staff. <i>name</i>	staff. <i>room</i>	department. <i>id</i>
1	Industrial Eng.	1	11	Bob	100	1
1	Industrial Eng.	1	21	Fran	101	1
2	Computer Sci.	2	20	Betsy	100	2
2	Computer Sci.	2	54	Pat	102	2
4	Chemistry	1	22	Frank	102	4
5	Physics	4	35	Sarah	200	5
7	Materials Sci.	5	40	Sam	10	7

# Summary of INNER JOINS

- Introduced INNER JOIN
  - `table1 INNER JOIN table2 ON table1.col1 = table2.col2`
  - Creates a virtual table
  - Rows are matched according to columns specified with “ON”
    - Usually this is a foreign key
    - If “ON” is omitted, all columns with identical names are checked for a match
  - Joined table has all the columns from both tables
- NOTE:
  - The “INNER” keyword is optional
  - If a matching row is not found in the second table, the row is omitted
  - In other words, a row must exist in both tables to produce a row in the joined table

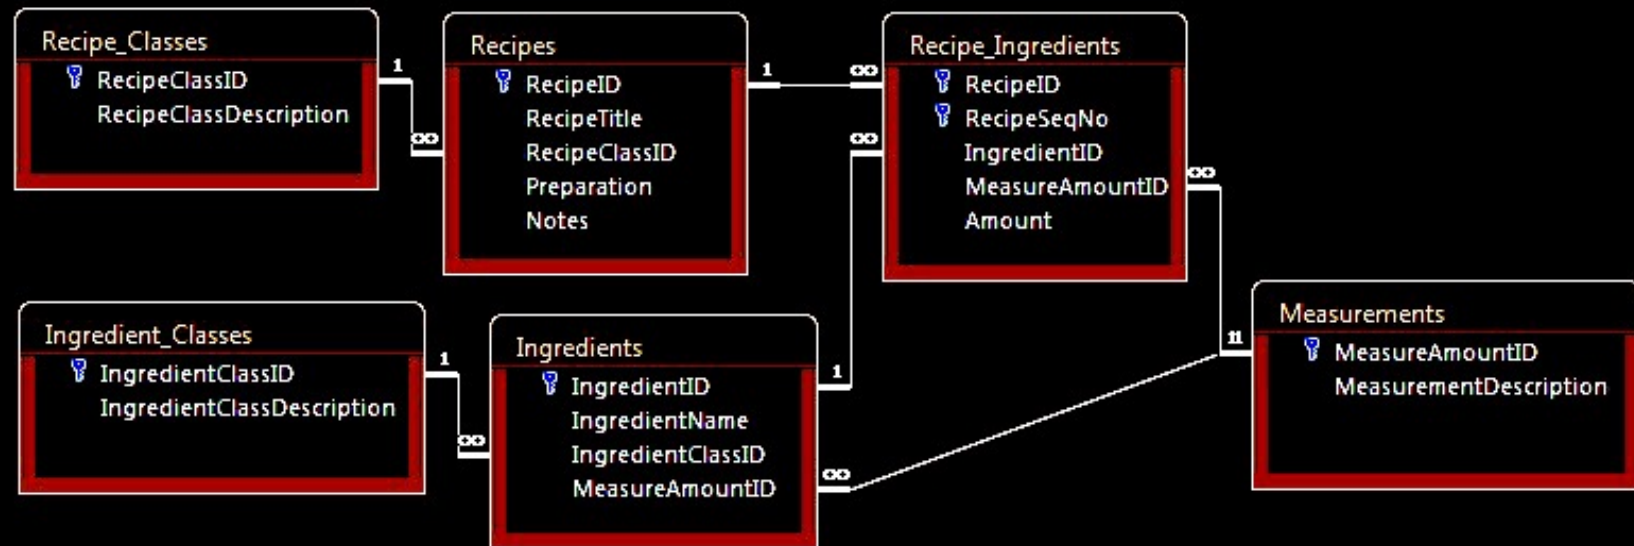
Print the recipe (ingredients, amount, measure) for Irish Stew in recipe sequence order (RecipeID = 1) (*Recipes.sqlite*)



Print the recipe (ingredients, amount, measure) for Irish Stew in recipe sequence order (RecipeID = 1) (*Recipes.sqlite*)

```
SELECT IngredientName,  
        Amount,  
        Measurements.MeasurementDescription  
FROM Recipe_Ingredients  
    JOIN Ingredients  
        ON Recipe_Ingredients.IngredientId  
           = Ingredients.IngredientID  
    JOIN Measurements  
        ON Recipe_Ingredients.MeasureAmountID  
           = Measurements.MeasureAmountID  
WHERE RecipeId=1  
ORDER BY RecipeSeqNo;
```

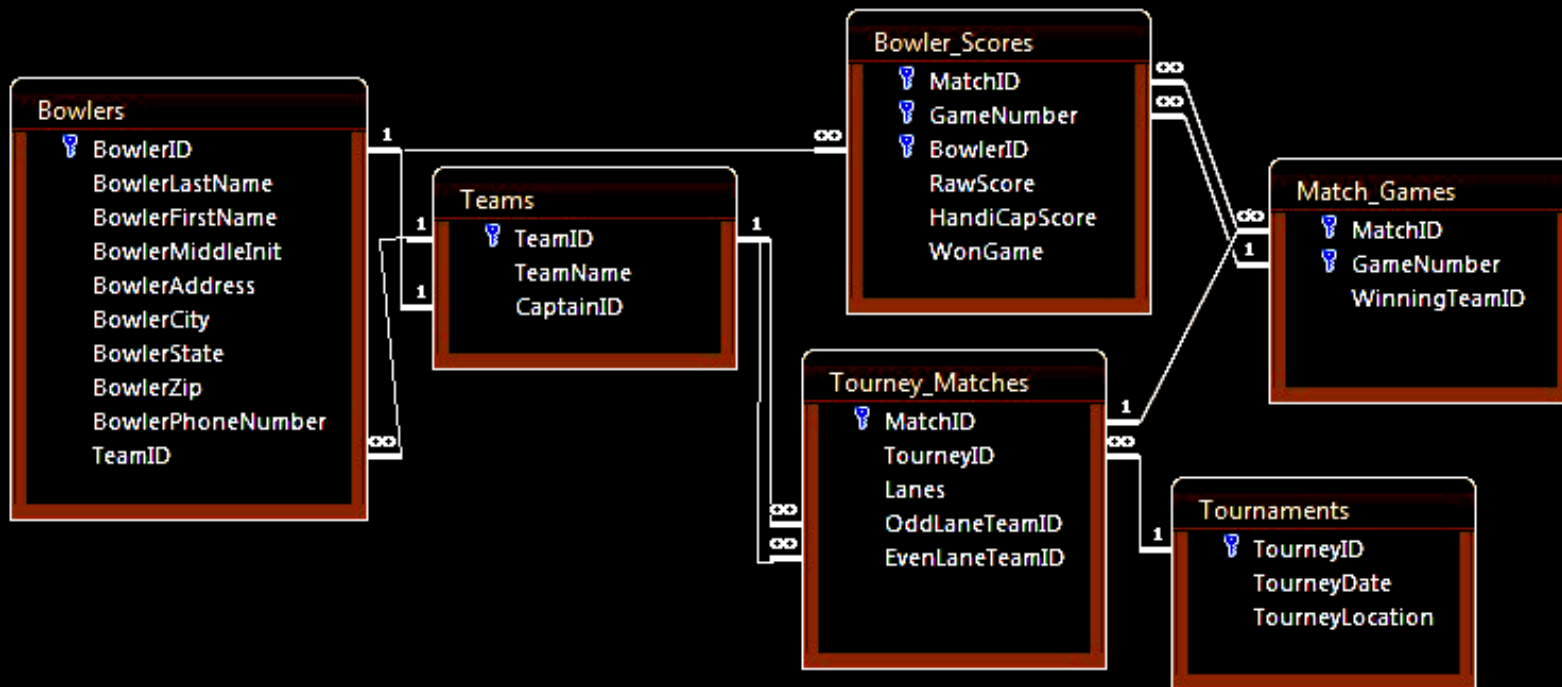
What is the name of the recipe with the most ingredients,  
and how many ingredients does it have? (*Recipes.sqlite*)  
(Can be done with either a subquery or a JOIN)



What is the name of the recipe with the most ingredients, and how many ingredients does it have? (*Recipes.sqlite*)

```
SELECT Recipes.RecipeTitle,  
       COUNT(DISTINCT IngredientID) AS NumIngredients  
FROM Recipe_Ingredients  
     JOIN Recipes  
       ON Recipes.RecipeID = Recipe_Ingredients.RecipeID  
GROUP BY Recipes.RecipeID  
ORDER BY NumIngredients DESC  
LIMIT 1;
```

Print a schedule of all the team matchups over the whole season (Date, Location, OddTeamName, EvenTeamName)  
(*BowlingLeague.sqlite*)




Print a schedule of all the team matchups over the whole season (Date, Location, OddTeamName, EvenTeamName)  
(*BowlingLeague.sqlite*)

```
SELECT TourneyDate,  
       TourneyLocation,  
       OddTeam.TeamName,  
       EvenTeam.TeamName  
FROM Tourney_Matches  
  JOIN Tournaments  
    ON Tourney_Matches.TourneyID = Tournaments.TourneyID  
  JOIN Teams AS OddTeam  
    ON OddLaneTeamID = OddTeam.TeamID  
  JOIN Teams AS EvenTeam  
    ON EvenLaneTeamID = EvenTeam.TeamID;
```



Print game results for Tournament #1, including bowler names, team names, & raw score (*BowlingLeague.sqlite*)



Print game results for Tournament #1, including bowler names, team names, & raw score (*BowlingLeague.sqlite*)

```
SELECT
    Bowler_Scores.MatchID, GameNumber, TeamName,
    BowlerFirstName || " " || BowlerLastName AS Bowler,
    RawScore
FROM Bowler_Scores
    JOIN Tourney_Matches
        ON Bowler_Scores.MatchID = Tourney_Matches.MatchID
    JOIN Bowlers
        ON Bowlers.BowlerID = Bowler_Scores.BowlerID
    JOIN Teams
        ON Bowlers.TeamID = Teams.TeamID
WHERE TourneyId=1;
```