

MSiA-413 Introduction to Databases and Information Retrieval

Lecture 11 NATURAL and LEFT JOINs

Instructor: Nikos Hardavellas

Slides adapted from Steve Tarzia, C. L. Moffatt

Last Lecture

Showed how to create and modify data in SQL databases

- `CREATE TABLE` defines columns, primary key, foreign keys
- `INSERT INTO` adds rows
 - Bulk loading
- `DELETE FROM` removes rows
- `UPDATE` changes column values for existing rows
- `ALTER TABLE` changes schema
- `CREATE INDEX` adds and index

NATURAL JOIN

- A shorthand notation to make some JOIN queries shorter to express
- NATURAL JOIN matches rows using **columns with identical names**

For example:

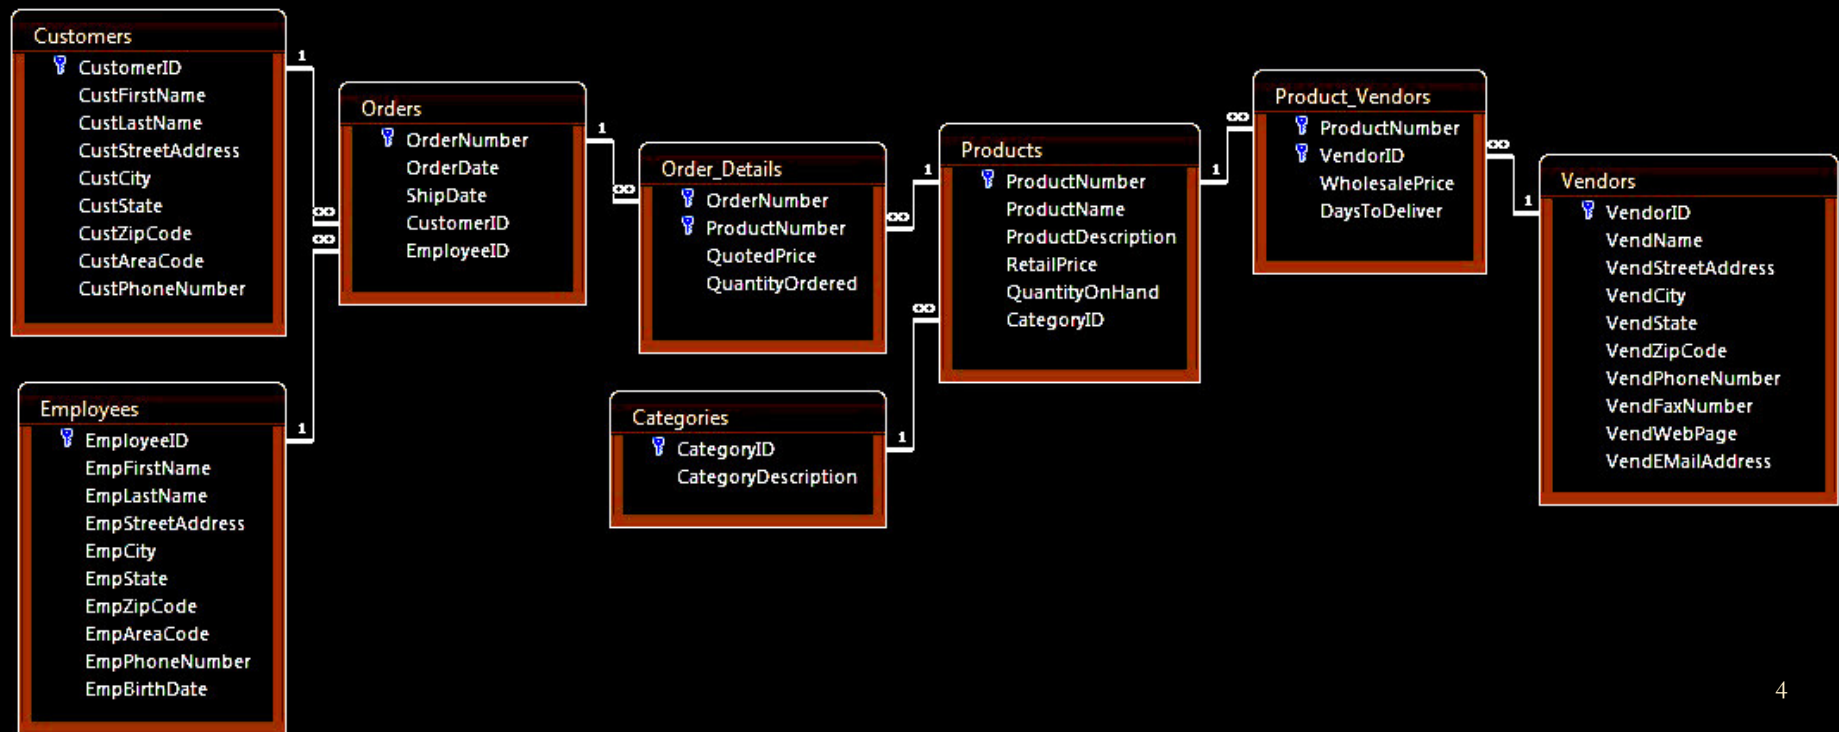
```
SELECT * FROM Orders JOIN Order_Details  
    ON Orders.OrderNumber = Order_Details.OrderNumber;
```

Becomes:

```
SELECT * FROM Orders NATURAL JOIN Order_Details;
```

Designing your data model NATURAL-ly


- Consistent column naming allows you to use NATURAL JOINS
- This is a good reason to avoid generic column names like “id” or “name”
 - Be explicit; use “CustomerID” and “EmployeeID”, etc, instead of just “id”




Different JOINS

- **INNER JOIN** constructs a table of all pairs of matching rows from two tables
 - **INNER** is the default
 - Useful for foreign keys
- However, there are many other ways to JOIN tables if you don't require matching


SQL JOINS




```
SELECT <select_list>  
FROM TableA A  
LEFT JOIN TableB B  
ON A.Key = B.Key
```




```
SELECT <select_list>  
FROM TableA A  
LEFT JOIN TableB B  
ON A.Key = B.Key  
WHERE B.Key IS NULL
```




```
SELECT <select_list>  
FROM TableA A  
INNER JOIN TableB B  
ON A.Key = B.Key
```




```
SELECT <select_list>  
FROM TableA A  
FULL OUTER JOIN TableB B  
ON A.Key = B.Key
```



```
SELECT <select_list>  
FROM TableA A  
RIGHT JOIN TableB B  
ON A.Key = B.Key
```



```
SELECT <select_list>  
FROM TableA A  
RIGHT JOIN TableB B  
ON A.Key = B.Key  
WHERE A.Key IS NULL
```




```
SELECT <select_list>  
FROM TableA A  
FULL OUTER JOIN TableB B  
ON A.Key = B.Key  
WHERE A.Key IS NULL  
OR B.Key IS NULL
```

© C.L. Moffatt, 2008

LEFT JOIN

- LEFT JOIN includes *all* rows in the first table (*left*-hand side) and just the matching rows in the second table (right-hand side)

LEFT JOIN



(standard)
INNER JOIN



LEFT JOIN output

- Like all JOINS, LEFT JOIN prints columns from the left table followed by columns from the right table
- However, with LEFT JOIN, some rows will have all *NULL* values in the right table columns, meaning that no match was found in the right table
- When to use LEFT JOIN?
 - To supplement a table with additional information that may be available for some rows, but not available for all the rows

LEFT JOIN example

staff			
<i>id</i>	<i>name</i>	<i>room</i>	<i>departmentId</i>
11	Bob	100	1
20	Betsy	100	NULL
21	Fran	101	1
22	Frank	102	99999
35	Sarah	200	5
40	Sam	10	7
54	Pat	102	2

department		
<i>id</i>	<i>name</i>	<i>buildingId</i>
1	Industrial Eng.	1
2	Computer Sci.	2
5	Physics	4
7	Materials Sci.	5

- Betsy and Frank have NULLs in the right half of the output because no matching department was found
- In other words, no pair of rows was found to satisfy the condition **ON** `staff.departmentId = department.id`

SELECT * FROM staff LEFT JOIN department ON staff.departmentId=department.id;

<i>staff.id</i>	<i>staff.name</i>	<i>staff.room</i>	<i>staff.departmentId</i>	<i>department.id</i>	<i>department.name</i>	<i>department.buildingId</i>
11	Bob	100	1	1	Industrial Eng.	1
20	Betsy	100	NULL	NULL	NULL	NULL
21	Fran	101	1	1	Industrial Eng.	1
22	Frank	102	99999	NULL	NULL	NULL
35	Sarah	200	5	5	Physics	4
40	Sam	10	7	7	Materials Sci.	5
54	Pat	102	2	2	Computer Sci.	2

staff			
<i>id</i>	<i>name</i>	<i>room</i>	<i>departmentId</i>
11	Bob	100	1
20	Betsy	100	NULL
21	Fran	101	1
22	Frank	102	99999
35	Sarah	200	5
40	Sam	10	7
54	Pat	102	2

LEFT JOIN example

department		
<i>id</i>	<i>name</i>	<i>buildingId</i>
1	Industrial Eng.	1
2	Computer Sci.	2
3	Physics	4
7	Materials Sci.	5



- Betsy and Frank have NULLs in the right half of the output because no matching department was found
- In other words, no pair of rows was found to satisfy the condition **ON** `staff.departmentId = department.id`

SELECT * FROM staff LEFT JOIN department ON staff.departmentId=department.id;

<i>staff.id</i>	<i>staff.name</i>	<i>staff.room</i>	<i>staff.departmentId</i>	<i>department.id</i>	<i>department.name</i>	<i>department.buildingId</i>
11	Bob	100	1	1	Industrial Eng.	1
20	Betsy	100	NULL	NULL	NULL	NULL
21	Fran	101	1	1	Industrial Eng.	1
22	Frank	102	99999	NULL	NULL	NULL
35	Sarah	200	5	5	Physics	4
40	Sam	10	7	7	Materials Sci.	5
54	Pat	102	2	2	Computer Sci.	2

LEFT JOIN example

staff			
<i>id</i>	<i>name</i>	<i>room</i>	<i>departmentId</i>
11	Bob	100	1
20	Betsy	100	NULL
21	Fran	101	1
22	Frank	102	99999
35	Sarah	200	5
40	Sam	10	7
54	Pat	102	2

department		
<i>id</i>	<i>name</i>	<i>buildingId</i>
1	Industrial Eng.	1
2	Computer Sci.	2
5	Physics	4
7	Materials Sci.	5

- Betsy and Frank have NULLs in the right half of the output because no matching department was found
- In other words, no pair of rows was found to satisfy the condition **ON** `staff.departmentId = department.id`

SELECT * FROM staff LEFT JOIN department ON staff.departmentId=department.id;

<i>staff.id</i>	<i>staff.name</i>	<i>staff.room</i>	<i>staff.departmentId</i>	<i>department.id</i>	<i>department.name</i>	<i>department.buildingId</i>
11	Bob	100	1	1	Industrial Eng.	1
20	Betsy	100	NULL	NULL	NULL	NULL
21	Fran	101	1	1	Industrial Eng.	1
22	Frank	102	99999	NULL	NULL	NULL
35	Sarah	200	5	5	Physics	4
40	Sam	10	7	7	Materials Sci.	5
54	Pat	102	2	2	Computer Sci.	2

LEFT JOIN example

staff			
<i>id</i>	<i>name</i>	<i>room</i>	<i>departmentId</i>
11	Bob	100	1
20	Betsy	100	NULL
21	Fran	101	1
22	Frank	102	99999
35	Sarah	200	5
40	Sam	10	7
54	Pat	102	2

department		
<i>id</i>	<i>name</i>	<i>buildingId</i>
1	Industrial Eng.	1
2	Computer Sci.	2
3	Physics	4
7	Materials Sci.	5



- Betsy and Frank have NULLs in the right half of the output because no matching department was found
- In other words, no pair of rows was found to satisfy the condition **ON** `staff.departmentId = department.id`

SELECT * FROM staff LEFT JOIN department ON staff.departmentId=department.id;

<i>staff.id</i>	<i>staff.name</i>	<i>staff.room</i>	<i>staff.departmentId</i>	<i>department.id</i>	<i>department.name</i>	<i>department.buildingId</i>
11	Bob	100	1	1	Industrial Eng.	1
20	Betsy	100	NULL	NULL	NULL	NULL
21	Fran	101	1	1	Industrial Eng.	1
22	Frank	102	99999	NULL	NULL	NULL
35	Sarah	200	5	5	Physics	4
40	Sam	10	7	7	Materials Sci.	5
54	Pat	102	2	2	Computer Sci.	2

LEFT JOIN with Grouping

- When computing an *aggregation* on a *many-to-one* relationship, LEFT JOIN includes rows from the parent table with no children
- In `SchoolScheduling.sqlite`, count the classes taught by each faculty member:
 - If you want this report to include faculty members teaching zero classes, you must use LEFT JOIN:

```
SELECT StaffID, SUM(ClassID IS NOT NULL) AS num_classes  
FROM Faculty NATURAL LEFT JOIN Faculty_Classes  
GROUP BY StaffID ORDER BY num_classes;
```

- The conditional “`ClassID IS NOT NULL`” returns “1” when true, “0” when false. The sum of ones and zeros per faculty counts the number of classes that faculty teaches
- Note that “`COUNT(*) AS num_classes`” would return “1” for faculty members with no classes, because the result table still includes every unmatched row from the left table

Aside: summing an indicator variable



- Two ways to count recipes with “fish” in description:
 - ```
SELECT COUNT(*)
FROM Recipes
WHERE RecipeTitle LIKE "%fish%";
```

    - WHERE clause keeps just the rows matching “fish,” then these rows are counted
  - ```
SELECT SUM(RecipeTitle LIKE "%fish%")  
FROM Recipes;
```


 - A column is created for every recipe indicating whether it matches “fish” or not
 - The column’s value will be **1** if it matches and **0** if not
 - Sum of all the ones and zeros will be the count of recipes with “fish” in description

RIGHT JOIN is symmetrical to LEFT

- Includes all rows from right table and matching rows from left table
- Reordering the tables makes a RIGHT JOIN a LEFT JOIN, so it is not necessary to use the RIGHT JOIN syntax



LEFT JOIN with exclusion




```
SELECT <select_list>
FROM TableA A
LEFT JOIN TableB B
ON A.Key = B.Key
WHERE B.Key IS NULL
```

- Includes rows from a table that *must not* match another table
- Useful for finding rows lacking something
- Just add a **WHERE** clause to look for *NULL* values in the right-hand side of the joined table
- For example, to determine which faculty members should be assigned a class:
 - ```
SELECT *
FROM Faculty NATURAL LEFT JOIN Faculty_Classes
WHERE ClassID IS NULL;
```
- Which classrooms are unused?
  - ```
SELECT *
FROM Class_Rooms NATURAL LEFT JOIN Classes
WHERE ClassID IS NULL;
```

Support Restrictions for JOINS in DB Engines

- FULL OUTER JOINS and RIGHT JOINS are not available in SQLite or MySQL
- You can *emulate* FULL OUTER JOIN with UNION (or UNION ... EXCEPT)



Note:

- FULL JOIN
 - OUTER JOIN
 - FULL OUTER JOIN
- ... all refer to the same thing


SalesOrders.sqlite: Which customers never ordered a helmet?

Solution 1 (using EXCEPT):

```
SELECT CustomerID FROM Customers
EXCEPT SELECT CustomerID FROM Customers
      NATURAL JOIN Orders
      NATURAL JOIN Order_Details
      NATURAL JOIN Products
WHERE ProductName LIKE "%Helmet%";
```

Solution 2 (using LEFT JOIN with exclusion):

```
SELECT CustomerID
FROM Customers
LEFT JOIN (SELECT DISTINCT CustomerID
          AS helmet_customer
          FROM Orders
          NATURAL JOIN Order_Details
          NATURAL JOIN Products
          WHERE ProductName LIKE "%Helmet%")
ON CustomerID = helmet_customer
WHERE helmet_customer IS NULL;
```



Recipes.sqlite: “Display missing types of recipes”

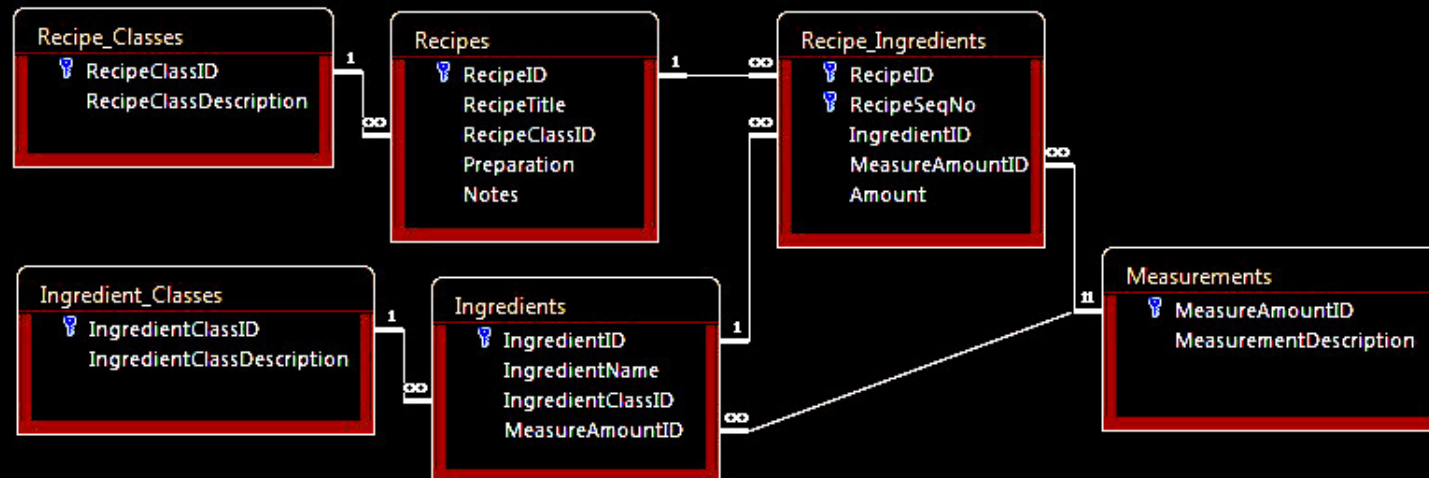
```
SELECT RecipeClassDescription  
FROM Recipe_Classes LEFT NATURAL JOIN Recipes  
GROUP BY RecipeClassID HAVING SUM(RecipeID IS NOT NULL) = 0;
```

or

```
SELECT RecipeClassDescription FROM Recipe_Classes  
WHERE RecipeClassID NOT IN (SELECT DISTINCT RecipeClassID FROM Recipes);
```

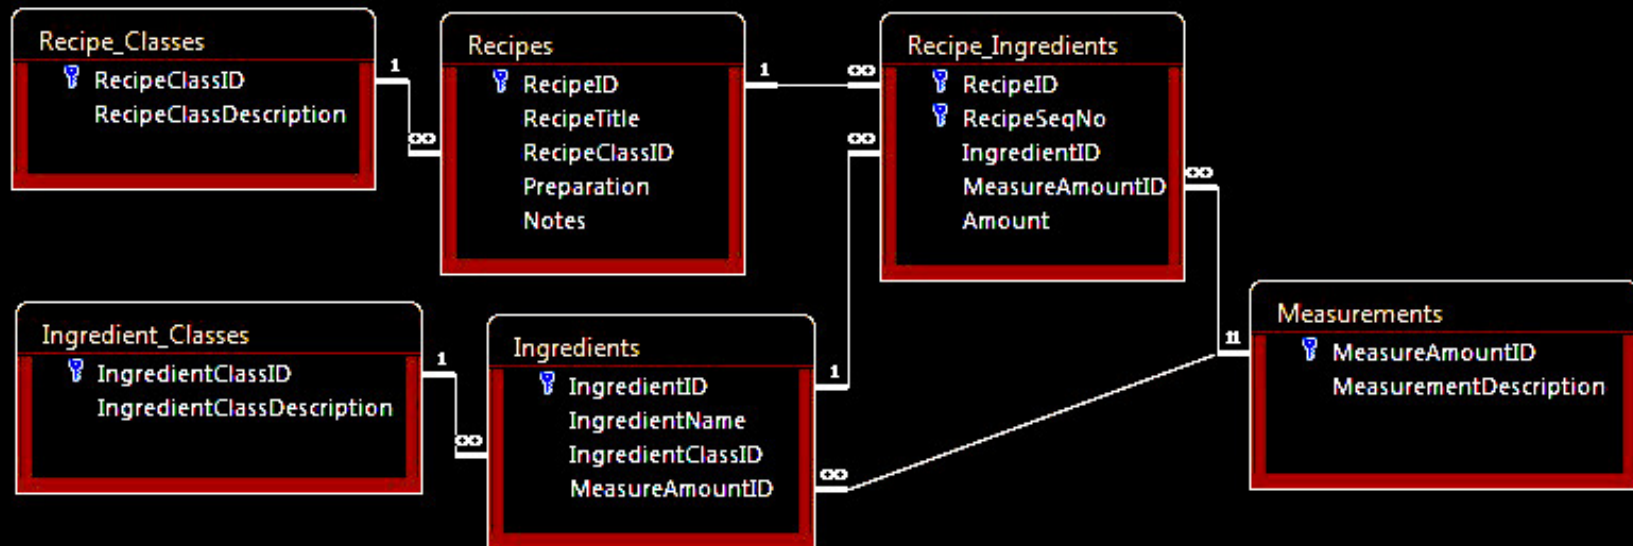
or

```
SELECT RecipeClassID FROM Recipe_Classes  
EXCEPT SELECT DISTINCT RecipeClassID FROM Recipes;
```



Recipes.sqlite: “List the number of recipes in each category (RecipeClassID)”

```
SELECT RecipeClassDescription,  
       SUM(RecipeID IS NOT NULL) AS RecipeCount  
FROM Recipe_Classes  
     LEFT NATURAL JOIN Recipes  
GROUP BY RecipeClassID;
```

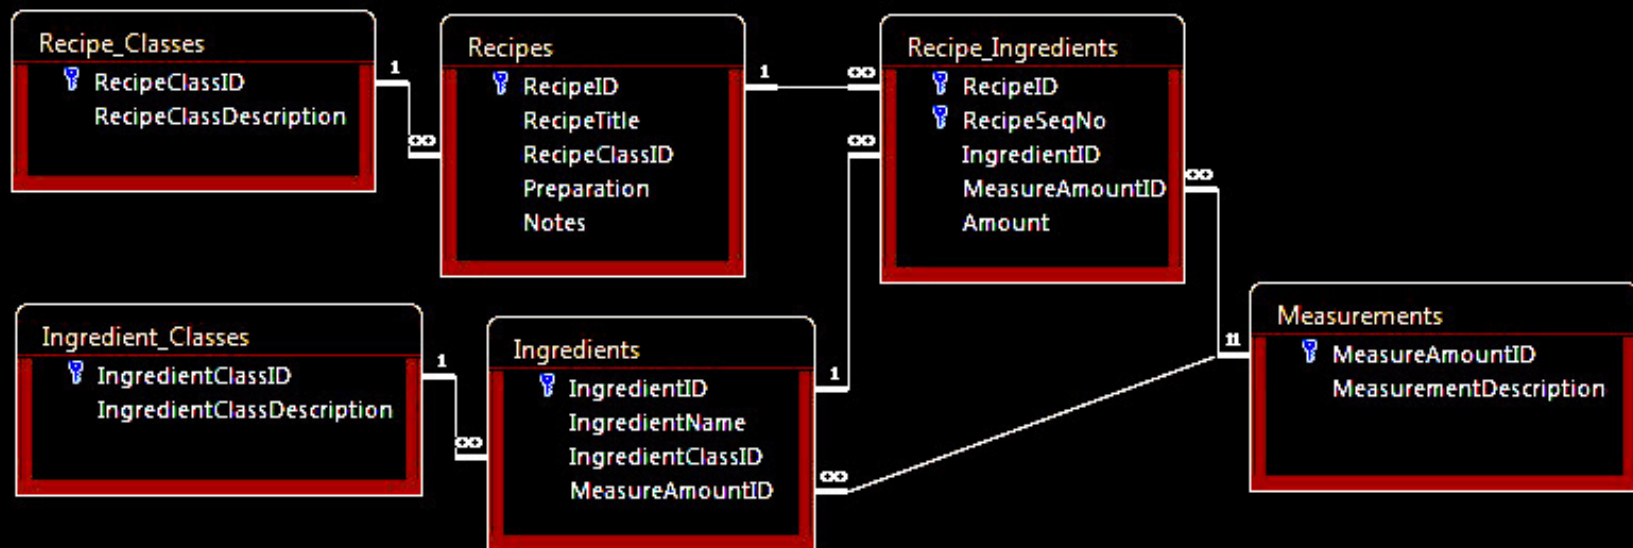


Recipes.sqlite: “Show me all ingredients and any recipes they are used in”

```
SELECT IngredientName, RecipeTitle FROM Ingredients  
LEFT NATURAL JOIN Recipe_Ingredients LEFT NATURAL JOIN Recipes;
```

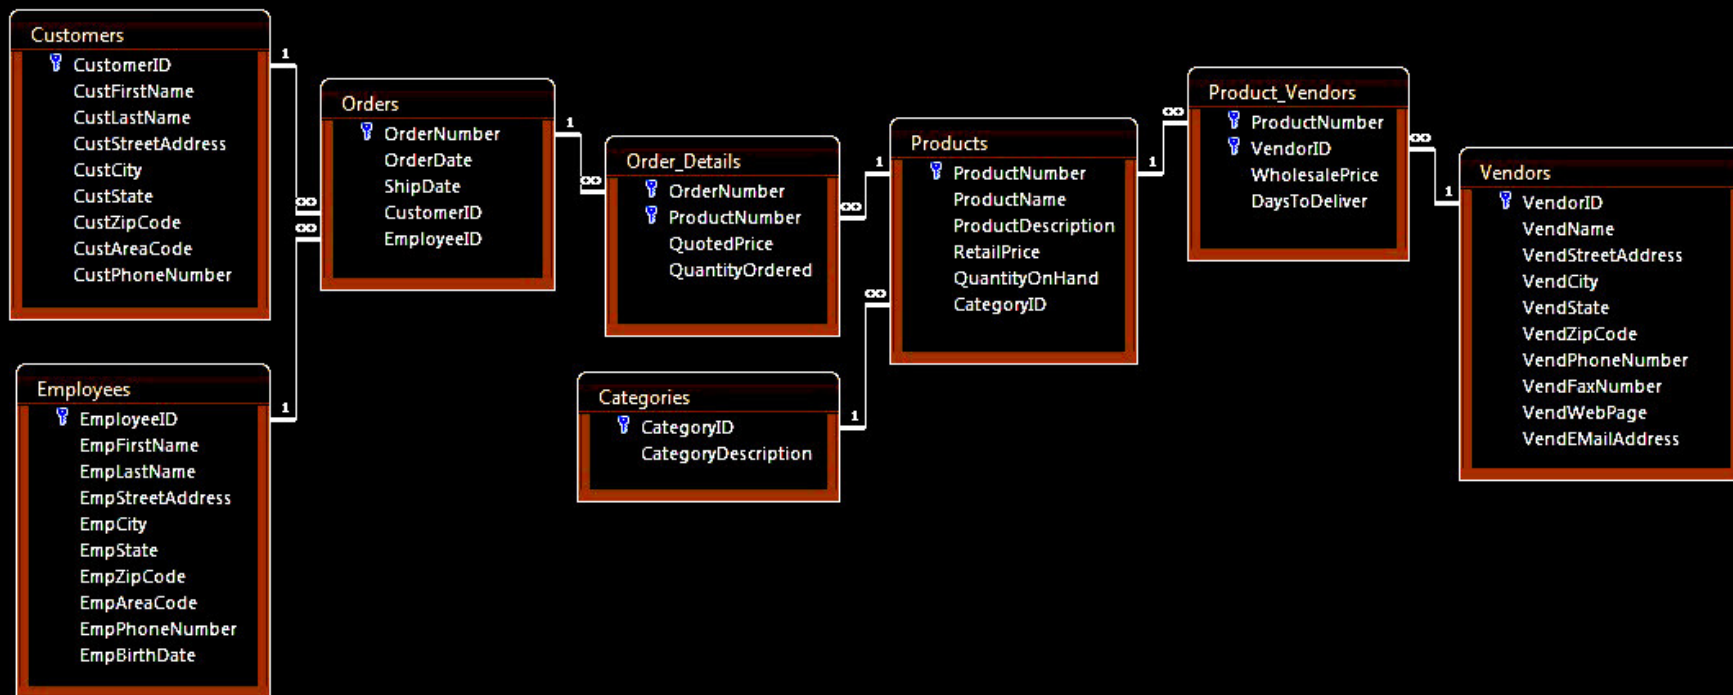
or, if only used ingredients are desired:

```
(Recipe_Ingredients NATURAL JOIN Recipes);
```



SalesOrders.sqlite: “Display customers who have no sales rep (employees) in the same ZIP Code”

- `SELECT * FROM Customers LEFT JOIN Employees ON CustZipCode=EmpZipCode WHERE EmpZipCode IS NULL;`
- `SELECT * FROM Customers WHERE CustZipCode IN (SELECT CustZipCode FROM Customers EXCEPT SELECT EmpZipCode FROM Employees);`



SalesOrders.sqlite: “List all products and the dates for any orders (of that product)”

```
SELECT Products.ProductNumber, ProductName, OrderDate
FROM Products
      LEFT NATURAL JOIN (Order_Details
                        NATURAL JOIN Orders);
```

