

MLDS-413 Introduction to Databases and Information Retrieval

Lecture 16 Window Functions Datetime Functions

Instructor: Nikos Hardavellas

Last Lecture


- Showed how indexes are added to tables
- Explained how multiple indexes can co-exist
- Described composite indexes
- Gave guidelines for when to index columns
- Showed which columns should be indexed to speed-up some example queries

SQL difficulties with aggregates


- Regular aggregation causes rows to become grouped into a single result row
- What if this is undesirable?
- Can you perform calculations for each row, based on nearby rows?
 - e.g., calculating moving average
- Window functions allow the use of a “window” of rows on calculations
 - This window is typically called a *frame*
 - The frame can be programmed as
 - a range
 - a partition
 - rows relative to the current row
 - ...
- Support for window functions was added to SQLite in 2018 (v3.25)

Window Functions

- Aggregate functions
 - Group rows into one
 - i.e., one result per group
 - Aggregate over entire group

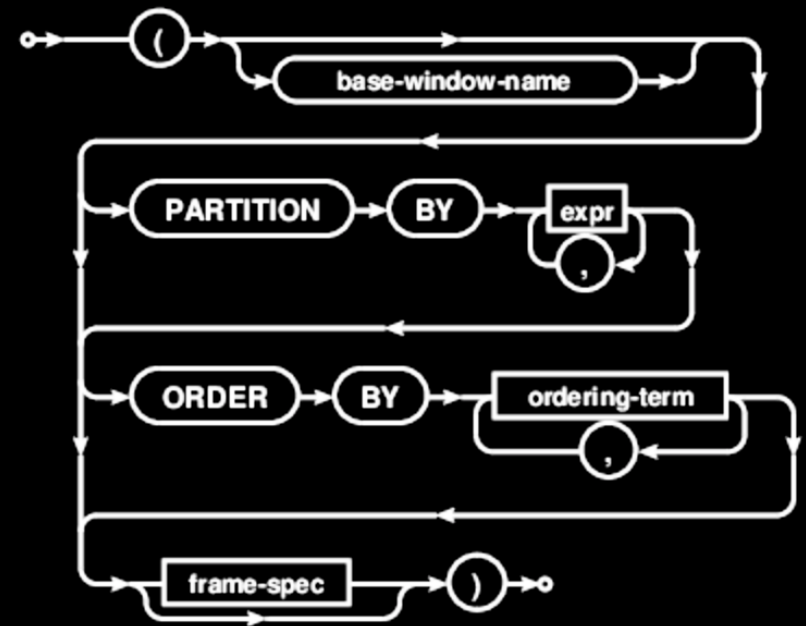


- Window functions
 - Retain row identity
 - i.e., one result per row
 - Aggregate over the frame
 - i.e., the “window” of rows



Window definitions

- Window partition (**PARTITION BY**)
 - Groups rows into partitions
 - Window-function processing is performed separately for each partition
 - Like a “*windowed group*”
- Window ordering (**ORDER BY**)
 - Defines the order or sequence of rows within each window
- Window frame (**ROWS**, **RANGE**, **GROUPS**)
 - Defines the window by use of an offset from the specified row, value or group



Window example: prepare GROUP BY #1

- Calculate the running total orders in `SalesOrder.sqlite` per month
 - i.e., if sales were \$10, \$10, \$30, the running total would be \$10, \$20, \$50
- First, let's calculate the total orders per individual month

```
SELECT OrderDate, SUM (OrderTotal)
FROM Orders
GROUP BY OrderDate;
```

	OrderDate	SUM (OrderTotal)
1	2012-09-01	52083.52
2	2012-09-02	56111.08
3	2012-09-03	34505.04
4	2012-09-04	23538.34



- ...this gives the total order per day, not **per month...**
- Correctness: need to convert date to month for **GROUP BY**
- Precision: it would be nice to print the last day of the month on the output
- We can do all these with **datetime functions**

Detour 1: Datetime functions: `timestring`

- A `timestring` is a character string in any of the following formats:
 - **YYYY-MM-DD HH:MM:SS.SSS[+-]HH:MM**
 - time is optional: HH:MM or HH:MM:SS or HH:MM:SS.SSS
 - time zone (following time) is also optional, or “**Z**” to indicate “Zulu” time (UTC)
 - ISO-8601 requires literal character “**T**” separating the date and the time
 - **HH:MM:SS.SSS**
 - Assume a date of 2000-01-01
 - seconds are optional: SS or SS.SSS
 - time zone is also optional
 - **now**
 - **DDDDDDDDDDDD**
 - Julian day, expressed as a floating-point value

Datetime functions

- `date(timestring, modifier, modifier, ...)` → YYYY-MM-DD
- `time(timestring, modifier, modifier, ...)` → HH:MM:SS
- `datetime(timestring, modifier, modifier, ...)` → YYYY-MM-DD HH:MM:SS
- `julianday(timestring, modifier, modifier, ...)` → Julian Day
 - The number of days since noon in Greenwich on November 24, 4714 B.C.
- `strftime(format, timestring, modifier, modifier, ...)` → date string
 - Date formatted according to the *format* string specified as the first argument:

%d	day of month: 00	%S	seconds since 1970-01-01
%f	fractional seconds: SS.SSS	%S	seconds: 00-59
%H	hour: 00-24	%w	day of week 0-6 with Sunday==0
%j	day of year: 001-366	%W	week of year: 00-53
%J	Julian day number	%Y	year: 0000-9999
%m	month: 01-12	%%	%
%M	minute: 00-59		

Datetime function modifiers

- The **timestring** can be followed by zero or more modifiers
- Modifiers precedence rules: transformations applied from left to right, in order
 - **[+ -]NN days**
 - **[+ -]NNN hours**
 - **[+ -]NNN minutes**
 - **[+ -]NNN.NNNN seconds**
 - **[+ -]NNN months**
 - Normalizes for months with <31 days
 - e.g., 2001-03-31 '+1 month'
→ 2001-04-31; April has only 30 days
→ normalize to 2001-05-01
 - **[+ -]NNN years**
 - Normalizes like above for leap years
 - **start of month**
 - **start of year**
 - **start of day**
 - **weekday N**
 - Advances date forward to desired weekday
 - Sunday = 0
 - **unixepoch**
 - Can only follow a DDDDDDDDDDD timestring
 - Causes timestring to be interpreted as Unix Epoch
 - **localtime**
 - Can only follow a timestring that is UTC
 - Causes timestring to be converted to local time
 - **utc**
 - Can only follow a timestring that is local time
 - Causes timestring to be converted to UTC

Datetime function examples

- Compute the current date
`SELECT date('now');`
- Compute the last day of the current month
`SELECT date('now', 'start of month', '+1 month', '-1 day');`
- Compute the date and time given a unix timestamp 1092941466, and translate to local time
`SELECT datetime(1092941466, 'unixepoch', 'localtime');`
- Compute the current unix timestamp (i.e., unix epoch time)
`SELECT strftime('%s', 'now');`
- Compute the number of days since the signing of the US Declaration of Independence
`SELECT julianday('now') - julianday('1776-07-04');`
- Compute the date of the first Tuesday in October for the current year
`SELECT date('now', 'start of year', '+9 months', 'weekday 2');`
- <https://www.sqlitetutorial.net/sqlite-date-functions/sqlite-date-function/>
- https://sqlite.org/lang_datefunc.html

Datetime function examples in PostgreSQL

- Compute the current date
`SELECT now(); --- YYYY-MM-DD HH:MM:SS in UTC`
`SELECT current_date; --- YYYY-MM-DD`
- Compute the first day of the current month
`SELECT date_trunc('month', current_date); --- YYYY-MM-01 00:00:00-04`
`SELECT date_trunc('month', current_date)::DATE; --- YYYY-MM-01`
- Compute the last day of the current month
`SELECT (date_trunc('month', current_date)
+ interval '1 month' - interval '1 day')::DATE;`
- Compute the date and time given a unix timestamp 1092941466, and translate to local time
`SELECT to_timestamp(1092941466);`
`SELECT to_timestamp(1092941466) AT TIME ZONE 'America/Chicago';`
- Compute the current unix timestamp (i.e., unix epoch time)
`SELECT extract(epoch FROM now());`
- Compute the number of days since the signing of the US Declaration of Independence
`SELECT current_date - '1776-07-04'::DATE;`
- <https://www.postgresql.org/docs/10/functions-datetime.html>

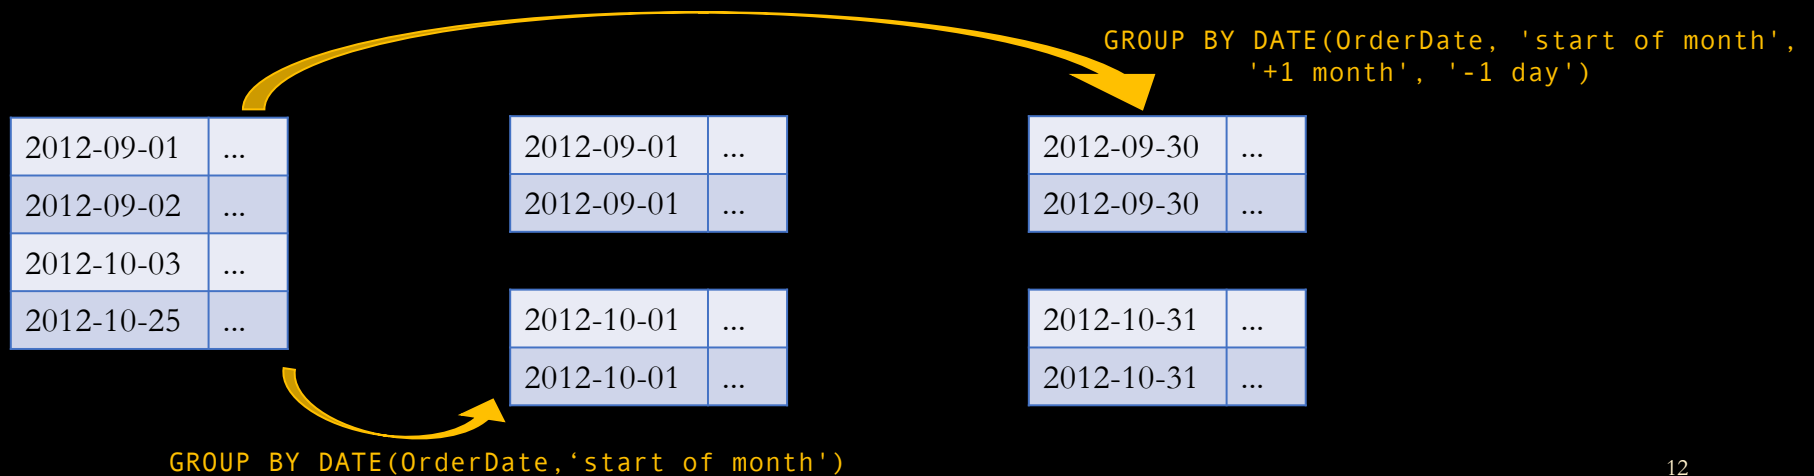
Window example: prepare GROUP BY #2

- Calculate the running total orders in `SalesOrder.sqlite` per month
- datetime → first day of the month

`DATE(OrderDate, 'start of month')`

- datetime → last day of the month

`DATE(OrderDate, 'start of month', '+1 month', '-1 day')`



Window example: prepare GROUP BY #3

- Calculate the running total orders in `SalesOrder.sqlite` per month
- datetime → first day of the month

`DATE(OrderDate, 'start of month')`

- datetime → last day of the month

`DATE(OrderDate, 'start of month', '+1 month', '-1 day')`

```
SELECT DATE(OrderDate, 'start of month', '+1 month', '-1 day') AS EndOfMonth,  
       SUM (OrderTotal) AS MonthlyTotal  
FROM Orders  
GROUP BY DATE(OrderDate, 'start of month');
```

- This query calculates the sum of orders **per month**

	EndOfMonth	MonthlyTotal
1	2012-09-30	820475.89
2	2012-10-31	726899.7700000001
3	2012-11-30	757151.65
4	2012-12-31	618910.4
5	2013-01-31	936516.5
6	2013-02-28	770941.16

Window example: prepare GROUP BY #4

- Calculate the running total orders in `SalesOrder.sqlite` per month
- First, let's turn the GROUP BY query into a CTE to use later:

```
WITH monthly_sales(EndOfMonth, MonthlyTotal) AS (  
    SELECT DATE(OrderDate, 'start of month', '+1 month', '-1 day'),  
           SUM (OrderTotal)  
    FROM Orders  
    GROUP BY DATE(OrderDate, 'start of month'))  
SELECT *  
FROM monthly_sales;
```

- Now we can use this to calculate the **running total**

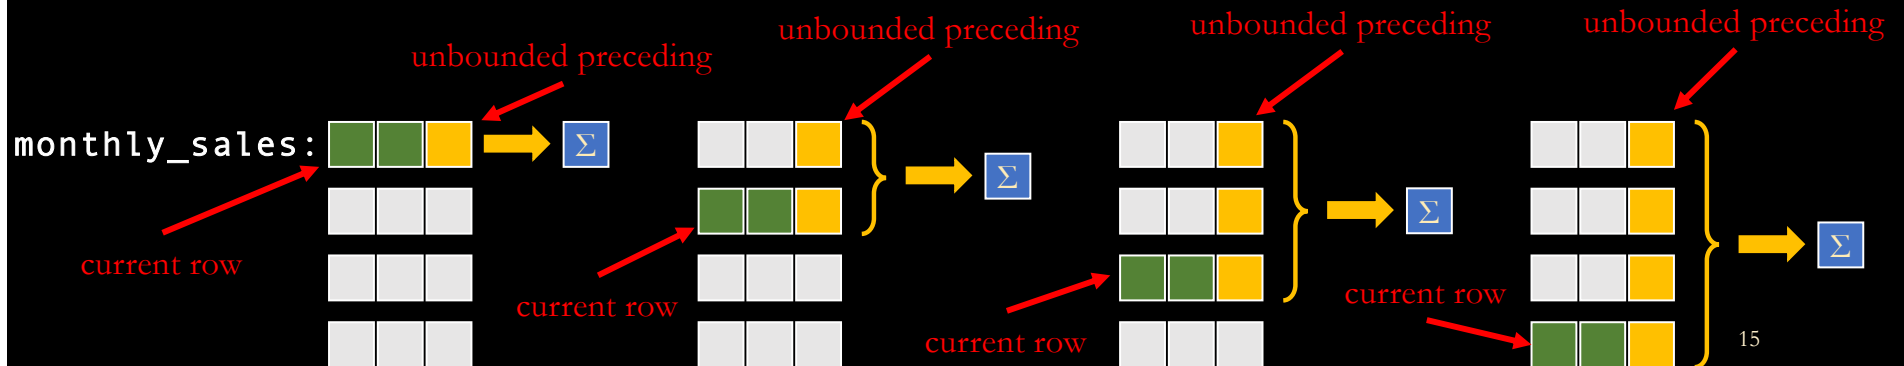
	EndOfMonth	MonthlyTotal
1	2012-09-30	820475.89
2	2012-10-31	726899.7700000001
3	2012-11-30	757151.65
4	2012-12-31	618910.4
5	2013-01-31	936516.5
6	2013-02-28	770941.16

Window example: running total orders per month

SalesOrder.sqlite

```
WITH monthly_sales(EndOfMonth, MonthlyTotal) AS (  
    SELECT DATE(OrderDate, 'start of month', '+1 month', '-1 day'),  
           SUM (OrderTotal)  
    FROM Orders  
    GROUP BY DATE(OrderDate, 'start of month'))  
SELECT EndOfMonth,  
       SUM(MonthlyTotal) OVER (ORDER BY EndOfMonth  
                               ROWS UNBOUNDED PRECEDING) AS RunningTotal  
FROM monthly_sales;
```

	EndOfMonth	RunningTotal
1	2012-09-30	820475.89
2	2012-10-31	1547375.66
3	2012-11-30	2304527.31
4	2012-12-31	2923437.71
5	2013-01-31	3859954.21
6	2013-02-28	4630895.37



ROWS vs. RANGE vs. GROUPS

- Consider a database with rows representing item quantities
- The windowed sums below will produce different results each

```
SELECT SequenceNum as RowNumber,  
       Quantity,  
       SUM(Quantity) OVER (ORDER BY Quantity ROWS 3 PRECEDING) AS SumOverRows,  
       SUM(Quantity) OVER (ORDER BY Quantity RANGE 3 PRECEDING) AS SumOverRange,  
       SUM(Quantity) OVER (ORDER BY Quantity GROUPS 3 PRECEDING) AS SumOverGroups  
FROM Items  
ORDER BY SequenceNum;
```


ROWS VS. RANGE VS. GROUPS

RowNumber	Quantity	SumOverRows	SumOverRange	SumOverGroups
1	1	1	2	2
2	1	2	2	2
3	2	5	5	5
4	5	10	13	15
5	5	14	13	15
6	7	20	17	22
7	8	25	49	52
8	8	28	49	52
9	8	31	49	52
10	8	32	49	52
11	9	33	48	58
12	10	35	58	58
13	11	38	84	84
14	11	41	84	84
15	11	43	84	84
16	22	55	22	74

current row - 3 = row 6

current row = row 9

Σ

```
SELECT
    SequenceNum as RowNumber,
    Quantity,
```

```
    SUM(Quantity) OVER (
        ORDER BY Quantity
        ROWS 3 PRECEDING)
        AS SumOverRows,
```

```
    SUM(Quantity) OVER (
        ORDER BY Quantity
        RANGE 3 PRECEDING)
        AS SumOverRange,
```

```
    SUM(Quantity) OVER (
        ORDER BY Quantity
        GROUPS 3 PRECEDING)
        AS SumOverGroups
```

```
FROM Items
ORDER BY SequenceNum;
```

ROWS VS. RANGE VS. GROUPS

RowNumber	Quantity	SumOverRows	SumOverRange	SumOverGroups
1	1	1	2	2
2	1	2	2	2
3	5	5	5	5
4	5	10	13	15
5	5	14	13	15
6	7	20	17	22
7	8	25	49	52
8	8	28	49	52
9	8	31	49	52
10	8	32	49	52
11	9	40	48	58
12	11	51	58	58
13	11	58	84	84
14	11	69	84	84
15	11	80	84	84
16	22	102	22	74

Annotations on the table:

- Yellow box: "current value - 3 = 5" with an arrow pointing to the Quantity value 5 in Row 3.
- Yellow box: "current row = row 9" with an arrow pointing to the RowNumber 9.
- Yellow box: "current value = 8" with an arrow pointing to the Quantity value 8 in Row 9.
- Red dashed box highlights the SumOverRange value 49 for Rows 7, 8, 9, and 10.
- A red bracket groups Rows 4 through 10, with a red arrow pointing to the SumOverRange value 49, accompanied by a red Σ symbol.

```

SELECT
    SequenceNum as RowNumber,
    Quantity,

    SUM(Quantity) OVER (
        ORDER BY Quantity
        ROWS 3 PRECEDING)
        AS SumOverRows,

    SUM(Quantity) OVER (
        ORDER BY Quantity
        RANGE 3 PRECEDING)
        AS SumOverRange,

    SUM(Quantity) OVER (
        ORDER BY Quantity
        GROUPS 3 PRECEDING)
        AS SumOverGroups

FROM Items
ORDER BY SequenceNum;
    
```

ROWS VS. RANGE VS. GROUPS

			overRows	SumOverRange	SumOverGroups
1	1		1	2	2
2	1		2	2	2
3	3		5	5	5
4	5		10	13	15
5	5		14	13	15
6	7		20	17	22
7	8		25	49	52
8	8		28	49	52
9	8		31	49	52
10	8		32	49	52
11	9		33	48	58
12	9		35	58	58
13	11		41	84	84
14	11		41	84	84
15	11		43	84	84
16	22		55	22	74

Annotations:

- current group - 3 (points to row 3)
- current row = row 9 (points to row 9)
- current group (points to row 9)
- Σ (points to row 7)

```

SELECT
  SequenceNum as RowNumber,
  Quantity,

  SUM(Quantity) OVER (
    ORDER BY Quantity
    ROWS 3 PRECEDING)
    AS SumOverRows,

  SUM(Quantity) OVER (
    ORDER BY Quantity
    RANGE 3 PRECEDING)
    AS SumOverRange,

  SUM(Quantity) OVER (
    ORDER BY Quantity
    GROUPS 3 PRECEDING)
    AS SumOverGroups

FROM Items
ORDER BY SequenceNum;

```

ROWS vs. RANGE vs. GROUPS full results

RowNumber	Quantity	SumOverRows	SumOverRange	SumOverGroups
1	1	1	2	2
2	1	2	2	2
3	3	5	5	5
4	5	10	13	15
5	5	14	13	15
6	7	20	17	22
7	8	25	49	52
8	8	28	49	52
9	8	31	49	52
10	8	32	49	52
11	9	33	48	58
12	10	35	58	58
13	11	38	84	84
14	11	41	84	84
15	11	43	84	84
16	22	55	22	74

```

SELECT
  SequenceNum as RowNumber,
  Quantity,

  SUM(Quantity) OVER (
    ORDER BY Quantity
    ROWS 3 PRECEDING)
    AS SumOverRows,

  SUM(Quantity) OVER (
    ORDER BY Quantity
    RANGE 3 PRECEDING)
    AS SumOverRange,

  SUM(Quantity) OVER (
    ORDER BY Quantity
    GROUPS 3 PRECEDING)
    AS SumOverGroups

FROM Items
ORDER BY SequenceNum;


```

Window example: running total orders per month

SalesOrder.sqlite

X SELECT

```
DATE(OrderDate, 'start of month', '+1 month', '-1 day') AS EndOfMonth,  
SUM(OrderTotal) OVER (  
    ORDER BY DATE(OrderDate, 'start of month')  
    GROUPS UNBOUNDED PRECEDING) AS RunningTotal  
FROM Orders;
```



	EndOfMonth	RunningTotal
1	2012-09-30	820475.89
2	2012-09-30	820475.89
3	2012-09-30	820475.89
⋮		
163	2012-09-30	820475.89
164	2012-10-31	1547375.66
165	2012-10-31	1547375.66



- ...but, there are many **duplicates** (one output per row)

Window example: running total orders per month

SalesOrder.sqlite

```
SELECT DISTINCT
```


```
    DATE(OrderDate, 'start of month', '+1 month', '-1 day') AS EndOfMonth,
```

```
    SUM(OrderTotal) OVER (
```

```
        ORDER BY DATE(OrderDate, 'start of month')
```

```
        GROUPS UNBOUNDED PRECEDING) AS RunningTotal
```

```
FROM Orders;
```



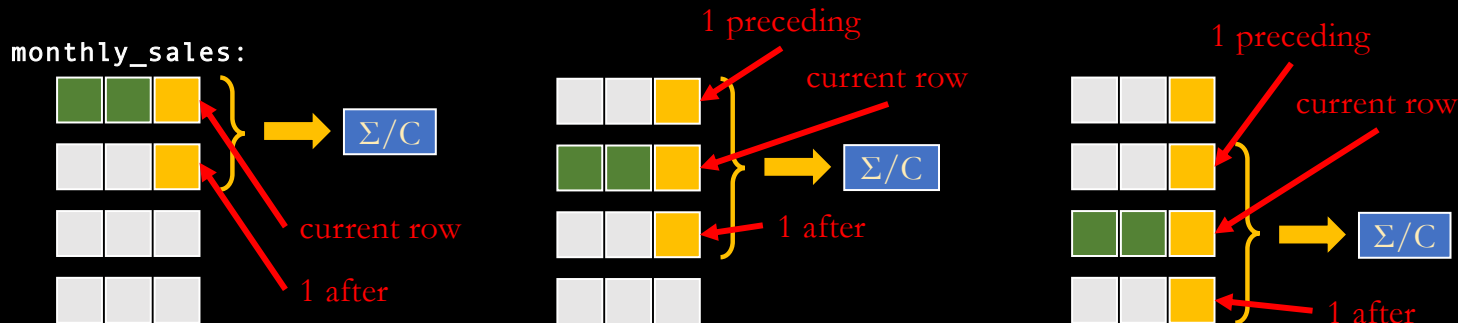
	EndOfMonth	RunningTotal
1	2012-09-30	820475.89
2	2012-10-31	1547375.66
3	2012-11-30	2304527.31
4	2012-12-31	2923437.71
5	2013-01-31	3859954.21
6	2013-02-28	4630895.37

Window example: running average of monthly orders

SalesOrder.sqlite

```
WITH monthly_sales(EndOfMonth, MonthlyTotal) AS (  
    SELECT DATE(OrderDate, 'start of month', '+1 month', '-1 day'),  
           SUM (OrderTotal)  
    FROM Orders  
    GROUP BY DATE(OrderDate, 'start of month'))  
SELECT EndOfMonth,  
       AVG(MonthlyTotal) OVER (ORDER BY EndOfMonth  
                               ROWS BETWEEN 1 PRECEDING AND 1 FOLLOWING)  
    AS RunningAvg  
FROM monthly_sales;
```

	EndOfMonth	RunningAvg
1	2012-09-30	773687.83
2	2012-10-31	768175.77
3	2012-11-30	700987.2733333333
4	2012-12-31	770859.5166666667
5	2013-01-31	775456.02
6	2013-02-28	853728.83

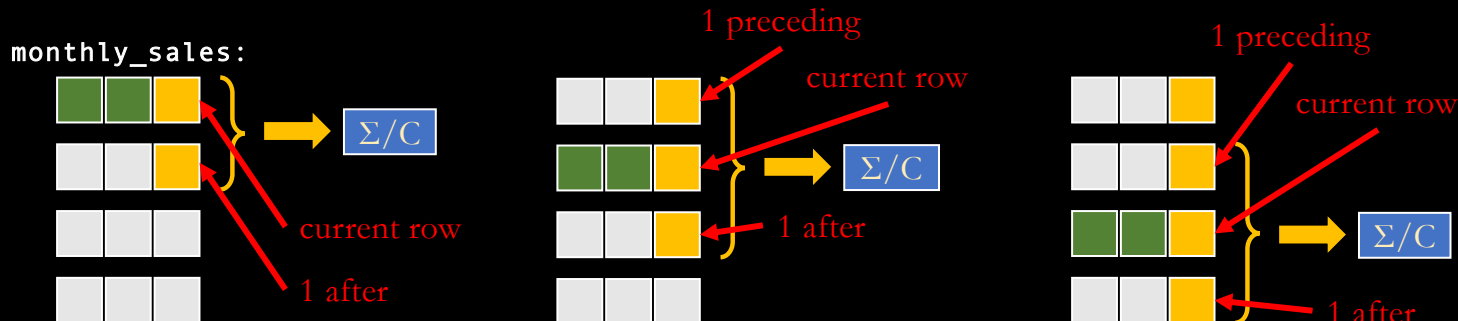


Window example: running average of monthly orders

SalesOrder.sqlite

```
WITH monthly_sales(EndOfMonth, MonthlyTotal) AS (  
    SELECT DATE(OrderDate, 'start of month', '+1 month', '-1 day'),  
           SUM (OrderTotal)  
    FROM Orders  
    GROUP BY DATE(OrderDate, 'start of month'))  
SELECT EndOfMonth,  
       AVG(MonthlyTotal) OVER win  
       AS RunningAvg  
FROM monthly_sales  
WINDOW win AS (ORDER BY EndOfMonth ROWS BETWEEN 1 PRECEDING AND 1 FOLLOWING);
```

	EndOfMonth	RunningAvg
1	2012-09-30	773687.83
2	2012-10-31	768175.77
3	2012-11-30	700987.2733333333
4	2012-12-31	770859.5166666667
5	2013-01-31	775456.02
6	2013-02-28	853728.83

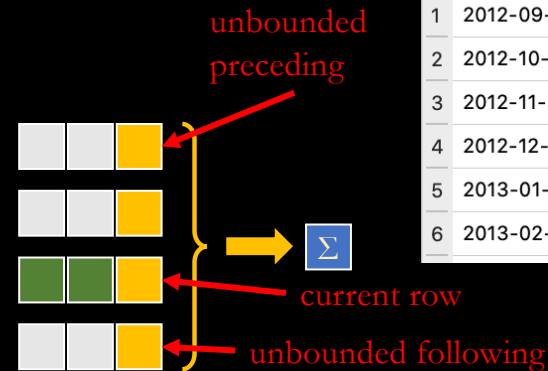
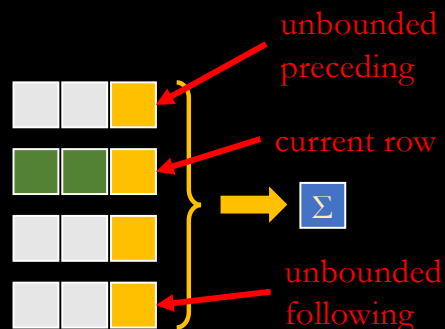



Window example: monthly orders as % of total

SalesOrder.sqlite

```
WITH monthly_sales(EndOfMonth, MonthlyTotal) AS (  
    SELECT DATE(OrderDate, 'start of month', '+1 month', '-1 day'),  
           SUM (OrderTotal)  
    FROM Orders  
    GROUP BY DATE(OrderDate, 'start of month'))  
SELECT EndOfMonth,  
       MonthlyTotal * 100.0 / SUM(MonthlyTotal) OVER (ORDER BY EndOfMonth  
                                                    ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING)  
       AS MonthlyPercent  
FROM monthly_sales;
```

monthly_sales:




	EndOfMonth	MonthlyPercent
1	2012-09-30	17.7174352786122
2	2012-10-31	15.6967435435709
3	2012-11-30	16.3500055497907
4	2012-12-31	13.3648107018233
5	2013-01-31	20.2232273712546
6	2013-02-28	16.6477775549483

Window example: average monthly shipping delay

SalesOrder.sqlite

- Solution using GROUP BY:

```
SELECT DATE(OrderDate, 'start of month', '+1 month', '-1 day') AS EndOfMonth,  
       AVG(julianday(ShipDate) - julianday(OrderDate))  
FROM Orders  
GROUP BY DATE(OrderDate, 'start of month');
```




	EndOfMonth	AvgDelay
1	2012-09-30	2.09815950920245
2	2012-10-31	2.08843537414966
3	2012-11-30	2.46206896551724
4	2012-12-31	2.2027972027972
5	2013-01-31	2.1027027027027
6	2013-02-28	2.59627329192547

Window example: average monthly shipping delay

SalesOrder.sqlite

- Solution using windows (partitions) attempt 1

```
X SELECT  
    DATE(OrderDate, 'start of month', '+1 month', '-1 day') AS EndOfMonth,  
    AVG(julianday(ShipDate) - julianday(OrderDate)) OVER win AS AvgDelay  
FROM Orders  
WINDOW win AS (PARTITION BY DATE(OrderDate, 'start of month'));
```



...but, there are many **duplicates** (one output per row)

	EndOfMonth	AvgDelay
1	2012-09-30	2.09815950920245
2	2012-09-30	2.09815950920245
3	2012-09-30	2.09815950920245
⋮		
163	2012-09-30	2.09815950920245
164	2012-10-31	2.08843537414966
165	2012-10-31	2.08843537414966




Window example: average monthly shipping delay

SalesOrder.sqlite

- Solution using windows (partitions) attempt 2

```
✓ SELECT DISTINCT  
    DATE(OrderDate, 'start of month', '+1 month', '-1 day') AS EndOfMonth,  
    AVG(julianday(ShipDate) - julianday(OrderDate)) OVER win AS AvgDelay  
FROM Orders  
WINDOW win AS (PARTITION BY DATE(OrderDate, 'start of month'));
```



	EndOfMonth	AvgDelay
1	2012-09-30	2.09815950920245
2	2012-10-31	2.08843537414966
3	2012-11-30	2.46206896551724
4	2012-12-31	2.2027972027972
5	2013-01-31	2.1027027027027
6	2013-02-28	2.59627329192547

Window Functions

Value	Ranking	Aggregate
FIRST_VALUE()	CUME_DIST()	AVG()
LAST_VALUE()	DENSE_RANK()	COUNT()
LAG()	NTILE()	MAX()
LEAD()	RANK()	MIN()
NTH_VALUE()	ROW_NUMBER()	SUM()
	PERCENT_RANK()	

Built-in Window Function Definitions

Value	Operation	Ranking	Operation
FIRST_VALUE()	Get the value of the first row in a specified window frame	CUME_DIST()	Compute the cumulative distribution of a value in an ordered set of values
LAST_VALUE()	Get the value of the last row in a specified window frame	DENSE_RANK()	Compute the rank for a row in an ordered set of rows with no gaps in rank values
LAG()	Provide access to a row at a given physical offset that comes before the current row	NTILE()	Divide a result set into a number of buckets as evenly as possible and assign a bucket number to each row
LEAD()	Provide access to a row at a given physical offset that follows the current row	PERCENT_RANK()	Calculate the percent rank of each row in an ordered set of rows
NTH_VALUE()	Return the value of an expression evaluated against the row N of the window frame in the result set	RANK()	Assign a rank to each row within the partition of the result set
		ROW_NUMBER()	Assign a sequential integer starting from one to each row within the current partition

Window ranking function example: ROW_NUMBER()


SalesOrder.sqlite

- Which customer issued the largest (in \$) order at each month?

```
WITH BigOrders(Month, CustomerID, OrderRank) AS (  
    SELECT strftime('%Y-%m', OrderDate), CustomerID,  
           ROW_NUMBER() OVER (PARTITION BY DATE(OrderDate, 'start of month')  
                               ORDER BY OrderTotal DESC) AS OrderRank  
    FROM Orders)  
SELECT Month, CustomerID,  
       CustFirstName || " " || CustLastName AS Name  
FROM BigOrders NATURAL JOIN Customers  
WHERE OrderRank = 1;
```

strftime('%Y-%m',
OrderDate)

	Month	CustomerID	Name
1	2012-09	1004	Robert Brown
2	2012-10	1017	Manuela Seidel
3	2012-11	1006	John Viescas
4	2012-12	1006	John Viescas
5	2013-01	1013	Rachel Patterson
6	2013-02	1005	Dean McCrae



Window Chaining

- A window can be defined by using another window as its base
- Implicitly copy the `PARTITION BY` and optionally `ORDER BY` clauses of the base window

```
SELECT b OVER (  
    win ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW  
)  
FROM t1  
WINDOW win AS (PARTITION BY a ORDER BY c);
```

- Rules of engagement:
 - No `PARTITION BY` clause in new window
 - No `ORDER BY` clause in new window if the base window has one
 - No frame specification in base window. The frame spec can only be given in the new window

Window Chaining Example

-- The following SELECT statement returns:

--

-- c | a | b | group_concat

-- one | 1 | A | A.D.G

-- one | 4 | D | D.G

-- one | 7 | G | G

-- three | 3 | C | C.F

-- three | 6 | F | F

-- two | 2 | B | B.E

-- two | 5 | E | E

--

```
SELECT c, a, b, group_concat(b, '.') OVER (  
  PARTITION BY c ORDER BY a RANGE BETWEEN CURRENT ROW AND UNBOUNDED FOLLOWING  
) AS group_concat  
FROM t1 ORDER BY c, a;
```

Compare to No Window Chaining

-- The following SELECT statement returns:

--

-- c | a | b | group_concat

-- one | 1 | A | A.D.G.C.F.B.E

-- one | 4 | D | D.G.C.F.B.E

-- one | 7 | G | G.C.F.B.E

-- three | 3 | C | C.F.B.E

-- three | 6 | F | F.B.E

-- two | 2 | B | B.E

-- two | 5 | E | E

--

```
SELECT c, a, b, group_concat(b, '.') OVER (  
  ORDER BY c, a ROWS BETWEEN CURRENT ROW AND UNBOUNDED FOLLOWING  
) AS group_concat  
FROM t1 ORDER BY c, a;
```

Window Chaining – SELECT's ORDER BY is immaterial

-- The following SELECT statement returns:

--

-- c | a | b | group_concat

-- one | 1 | A | A.D.G

-- two | 2 | B | B.E

-- three | 3 | C | C.F

-- one | 4 | D | D.G

-- two | 5 | E | E

-- three | 6 | F | F

-- one | 7 | G | G

--

```
SELECT c, a, b, group_concat(b, '.') OVER (
  PARTITION BY c ORDER BY a RANGE BETWEEN CURRENT ROW AND UNBOUNDED FOLLOWING
) AS group_concat
FROM t1 ORDER BY a;
```

Window Filters

FILTER clause: 

- If a filter clause is specified
 - Input rows are evaluated against the filter clause
 - If true → the row is fed to the window function
 - If false → the row is discarded
 - Can only be used with aggregate window functions (no value, ranking)

```
SELECT COUNT(*) FILTER (WHERE Quantity > 10)
      AS FilteredCount
FROM Items;
```

FilteredCount	
1	4

```
SELECT COUNT(*) FILTER (WHERE Quantity > 10)
      AS FilteredCount
FROM Items WHERE Quantity < 12;
```

FilteredCount	
1	3

RowNumber	Quantity
1	1
2	1
3	3
4	5
5	5
6	7
7	8
8	8
9	8
10	8
11	9
12	10
13	11
14	11
15	11
16	22

FILTER \neq WHERE

example_table:

	value	category
1	10	A
2	NULL	A
3	20	B
4	30	B

```
SELECT category,  
       COUNT(value)  
       AS count_values  
FROM example_table  
WHERE value > 15  
GROUP BY category;
```

	category	count_values
1	B	2

example_table:

	value	category
1	10	A
2	NULL	A
3	20	B
4	30	B


```
SELECT category,  
       COUNT(*)  
       AS count_values  
FROM example_table  
FILTER(WHERE value > 15)  
GROUP BY category;
```

	category	count_values
1	A	0
2	B	2

Frame Exclusion


- **EXCLUDE CURRENT ROW**

- Remove the current row itself from the frame
- Independent of the frame unit—remove even if the **RANGE** or **GROUPS** unit is used, and the current row has peers




- **EXCLUDE GROUP**

- Remove the current row + all its peers from the frame
- Independent of the frame unit



- **EXCLUDE TIES**

- Remove the peers of the current row, but not the current row itself

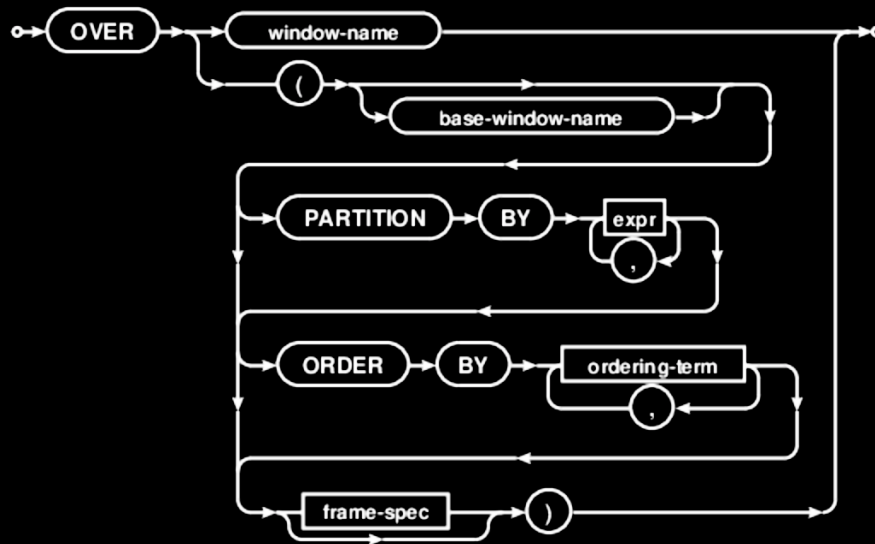


- **EXCLUDE NO OTHERS**

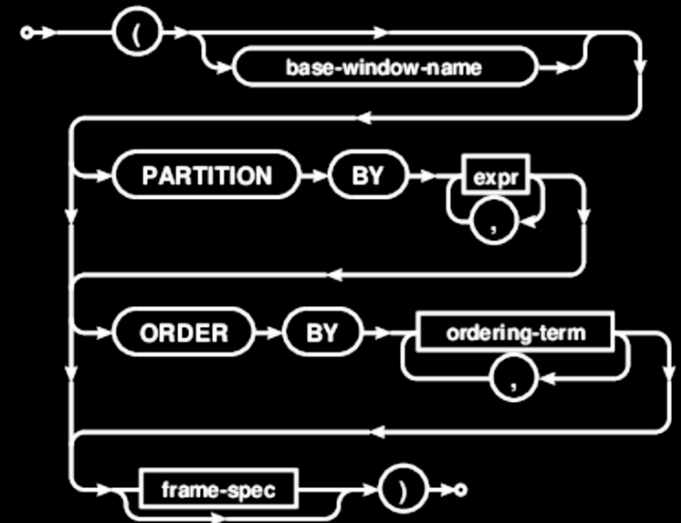
- does not remove any rows

Window Syntax Diagrams


OVER clause:



Window definition:



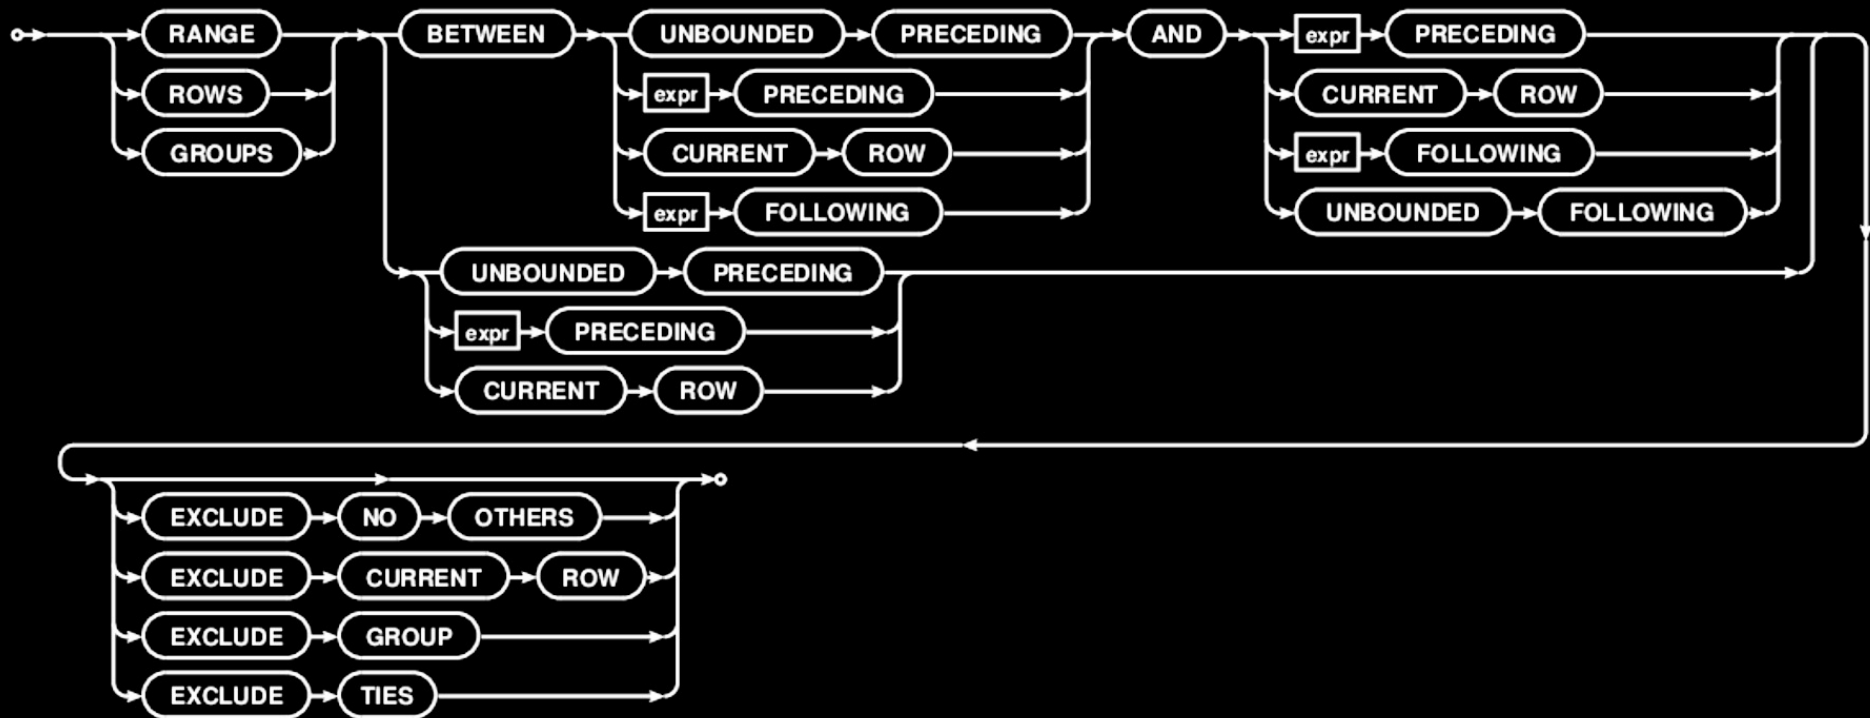
Window function invocation:



FILTER clause:

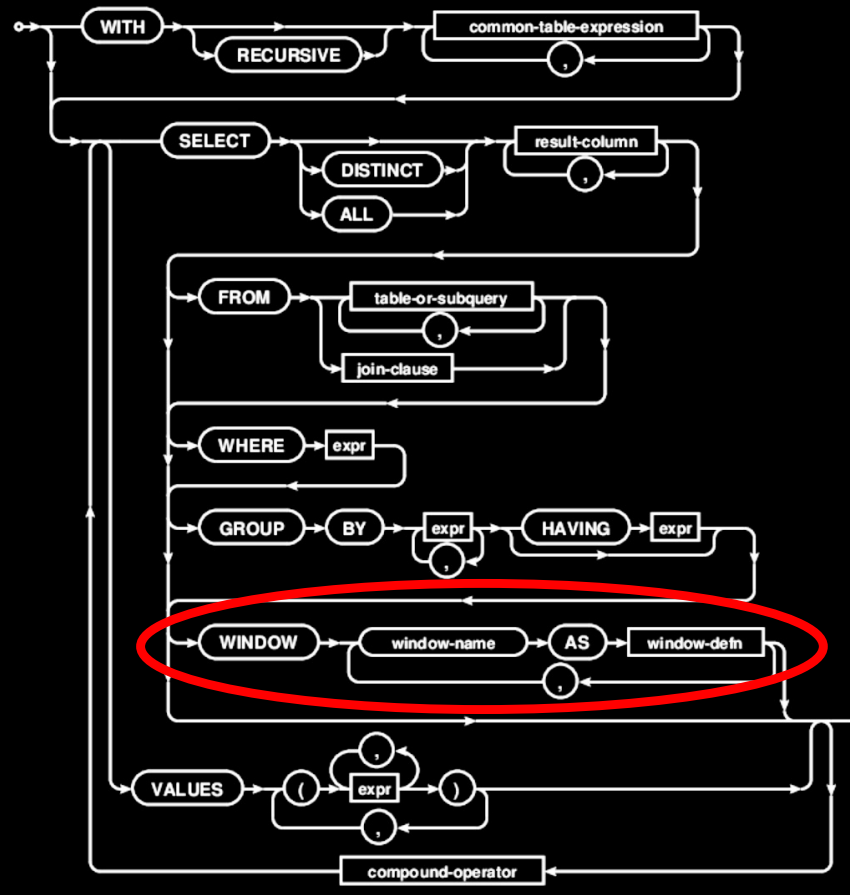
Window Syntax Diagrams

Frame specification:



Window Syntax Diagrams

Windowing in SELECT:



Windows in PostgreSQL

- Window frames with **ROWS**: supported since 2009
- Window frames with **RANGE**: full support added in 2018 with PostgreSQL 11
 - Previously only **UNBOUNDED PRECEDING** or **CURRENT ROW**
- Window frames with **GROUPS**: full support added in 2018 with PostgreSQL 11
 - Previously not supported *at all*
- **EXCLUDE**: support added in 2018 with PostgreSQL 11
 - Previously not supported *at all*
- More details at <https://www.postgresql.org/docs/10/functions-window.html>