

MLDS-413 Introduction to Databases and Information Retrieval

Lecture 17 Triggers Introduction to Transactions

Instructor: Nikos Hardavellas

Slides adapted from S. Tarzia, A. Silberschatz, H.F. Korth, S. Sudarshan

Last Lecture

- Datetime functions
 - Powerful functions to manipulate date and time
- Windowing
 - Constrain aggregators to a moving window of rows
- **OVER** and **WINDOW** statements
 - Define a window partition, order, and frame

Adding integrity constraints

SalesOrders.sqlite

- *New rule: no more than 4 employees in the same office (i.e., with same area code)*

```
SELECT EmpAreaCode, COUNT(*) AS NumEmployeesAtOffice
FROM Employees
GROUP BY EmpAreaCode;
```

EmpAreaCode	NumEmployeesAtOffice
206	1
210	1
253	2
425	4
515	1

- Inserting new employee at area code 425 should fail

```
INSERT INTO Employees
VALUES (800, "Nikos", "Hardavellas", "2233 Tech Dr",
      "Evanston", "IL", "60208", 425, "491-2270");
```

How to enforce the rule? Triggers

- A **trigger** is a statement that is executed automatically by the system as a side effect of a modification to the database
- To design a trigger mechanism, we must:
 - Specify the conditions under which the trigger is to be executed
 - Specify the actions to be taken when the trigger executes
- Complicated syntax
 - DB Browser for SQLite comment on github bug report, Aug 20, 2018:
“We practically don’t handle triggers at all in our application because they are complicated to parse”
 - Support is there now

Adding a trigger

SalesOrders.sqlite

- *New rule: no more than 4 employees in the same office (i.e., with same area code)*

```
CREATE TRIGGER Max4EmployeesPerOffice
BEFORE INSERT
ON Employees
FOR EACH ROW
BEGIN
    SELECT CASE
        WHEN (SELECT COUNT(*)
              FROM Employees
              WHERE EmpAreaCode = new.EmpAreaCode) >= 4
        THEN RAISE(FAIL, "Error: max 4 employees per office")
        END;
END;
```

Trigger name

When to "fire"

On which table

Run trigger per row

Refer to "new row"

Trigger actions

Raise an exception & print error

Insert with the trigger defined

SalesOrders.sqlite

- Inserting new employee at area code 425 should fail

```
INSERT INTO Employees  
VALUES (800, "Nikos", "Hardavellas", "2233 Tech Dr",  
        "Evanston", "IL", "60208", 425, "491-2270");
```

Execution finished with errors.

Result: Error: max 4 employees per office

At line 5:

```
INSERT INTO Employees  
VALUES (800, "Nikos", "Hardavellas", "2233 Tech Dr",  
        "Evanston", "IL", "60208", 425, "491-2270");
```

Trigger Events

- The trigger event can be an **insert**, **delete**, **update**, or **update of <cols>**

```
CREATE TRIGGER trigger_name  
BEFORE INSERT ON table  
...
```

```
CREATE TRIGGER trigger_name  
BEFORE DELETE ON table  
...
```

```
CREATE TRIGGER trigger_name  
BEFORE UPDATE ON table  
...
```

```
CREATE TRIGGER trigger_name  
BEFORE UPDATE OF column1, column2, ... ON table  
...
```

Trigger Timing

- The trigger can fire **before**, **after**, or **instead of** the triggering event

```
CREATE TRIGGER trigger_name  
BEFORE INSERT ON table  
...
```

← Typical use: add
integrity constraints

```
CREATE TRIGGER trigger_name  
INSERT ON table  
...
```

← Default is **BEFORE**

```
CREATE TRIGGER trigger_name  
AFTER INSERT ON table  
...
```

← Typical use: perform
additional actions

```
CREATE TRIGGER trigger_name  
INSTEAD OF INSERT ON table  
...
```


INSTEAD OF Triggers

- Changes the statement to execute
- Example: instead of modifying a view, modify the main table

```
CREATE VIEW CustomerView  
AS SELECT CustomerID, CustAreaCode FROM Customers;
```

```
SELECT CustAreaCode FROM customerview WHERE CustomerID = 1001;  
425
```

```
UPDATE CustomerView SET CustAreaCode = 314  
WHERE CustomerID = 1001;
```

Execution finished with errors.

Result: cannot modify CustomerView because it is a view

At line 18:

```
UPDATE CustomerView SET CustAreaCode = 314  
WHERE CustomerID = 1001
```

INSTEAD OF Triggers

- Changes the statement to execute
- Example: instead of modifying a view, modify the main table

```
CREATE VIEW CustomerView  
AS SELECT CustomerID, CustAreaCode FROM Customers;
```

```
CREATE TRIGGER CustomerViewChange  
INSTEAD OF UPDATE OF CustAreaCode ON CustomerView  
BEGIN  
    UPDATE customers SET CustAreaCode = new.CustAreaCode  
    WHERE CustomerID = new.CustomerID;  
END;
```

```
UPDATE CustomerView SET CustAreaCode = 314  
WHERE CustomerID = 1001;  
SELECT CustAreaCode FROM Customers WHERE CustomerID = 1001;  
314
```

Referencing attributes of old/new rows

- Example: data integrity in order \$ between Orders and Order_Details
- Without a trigger

```
SELECT OrderTotal FROM Orders WHERE OrderNumber=522;  
4.99
```

```
SELECT QuotedPrice FROM Order_Details WHERE OrderNumber=522;  
4.99
```

```
UPDATE Order_Details SET QuotedPrice=5.99  
WHERE OrderNumber=522;
```

```
SELECT OrderTotal FROM Orders WHERE OrderNumber=522;  
4.99
```

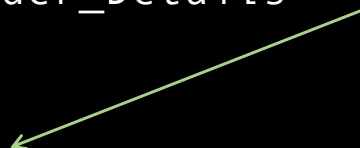
```
SELECT QuotedPrice FROM Order_Details WHERE OrderNumber=522;  
5.99
```

Referencing attributes of old/new rows

- Use **old.** and **new.** to refer to the old/new rows of the insert, update, or delete statement that fired the trigger
- **INSERT**: **new.** references are valid
- **UPDATE**: **new.** and **old.** references are valid
- **DELETE**: **old.** references are valid
- Example trigger for Orders and Order_Details

```
CREATE TRIGGER PriceChange
AFTER UPDATE OF QuotedPrice ON Order_Details
BEGIN
    UPDATE Orders
    SET OrderTotal = OrderTotal
      + old.QuantityOrdered * (new.QuotedPrice - old.QuotedPrice)
    WHERE OrderNumber = old.OrderNumber;
END;
```

new refers to the row
after the update



old refers to the row before the update



Referencing attributes of old/new rows

- Example: data integrity in order \$ between Orders and Order_Details
- With the trigger

```
SELECT OrderTotal FROM Orders WHERE OrderNumber=522;  
4.99
```

```
SELECT QuotedPrice FROM Order_Details WHERE OrderNumber=522;  
4.99
```

```
UPDATE Order_Details SET QuotedPrice=5.99  
WHERE OrderNumber=522;
```

```
SELECT OrderTotal FROM Orders WHERE OrderNumber=522;  
5.99
```

```
SELECT QuotedPrice FROM Order_Details WHERE OrderNumber=522;  
5.99
```


Trigger execution granularity

- Defines how often the trigger will execute

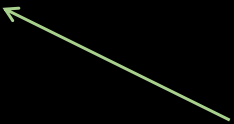
- Example:

```
CREATE TRIGGER TriggerName  
BEFORE INSERT  
ON table  
[FOR EACH ROW | FOR EACH STATEMENT]  
BEGIN  
...  
END;
```

Execute trigger once for
each row inserted by the
triggering statement



Execute trigger once for
each triggering statement



- SQLite implements only per-row triggers, hence this clause is optional

Trigger event filtering

- Execute trigger only when certain conditions are satisfied
- WHEN condition can access **new.** and **old.** rows
- Example: the previous trigger for Orders and Order_Details can be modified to detect updates of QuotedPrice

```
CREATE TRIGGER PriceChange
AFTER UPDATE OF QuotedPrice ON Order_Details
WHEN old.QuotedPrice <> new.QuotedPrice
BEGIN
    UPDATE Orders
    SET OrderTotal = OrderTotal
        + old.QuantityOrdered * (new.QuotedPrice -
old.QuotedPrice)
    WHERE OrderNumber = old.OrderNumber;
END;
```

Undefined behavior and **BEFORE** triggers

- Rules for **BEFORE UPDATE** and **BEFORE DELETE** triggers
 - If trigger modifies or deletes a row that was to have been updated or deleted
→ the subsequent update or delete operation is undefined
 - If trigger modifies or deletes a row
→ **AFTER** triggers that would have otherwise run on those rows may/may not run
- Rules for **BEFORE INSERT** triggers
 - If `rowid` is not explicitly set to an integer
→ **NEW.rowid** is undefined
- Because of the behaviors described above, programmers are encouraged to prefer **AFTER** triggers over **BEFORE** triggers when the triggers change data

Raising exceptions

- Notify the caller that an error has occurred
 - Actions: print an error message, return an error to the application if needed
- **RAISE()** is a **function** (i.e., part of an **expression**), not a statement
 - Must be within a **SELECT**, **CASE**, or any other statement accepting expressions

```
CREATE TRIGGER Max4EmployeesPerOffice
BEFORE INSERT ON Employees
FOR EACH ROW
BEGIN
    SELECT CASE
        WHEN (SELECT COUNT(*)
              FROM Employees
              WHERE EmpAreaCode = new.EmpAreaCode) >= 4
        THEN RAISE(FAIL, "Error: max 4 employees per office")
        END;
```


END;

Conflict resolution algorithm

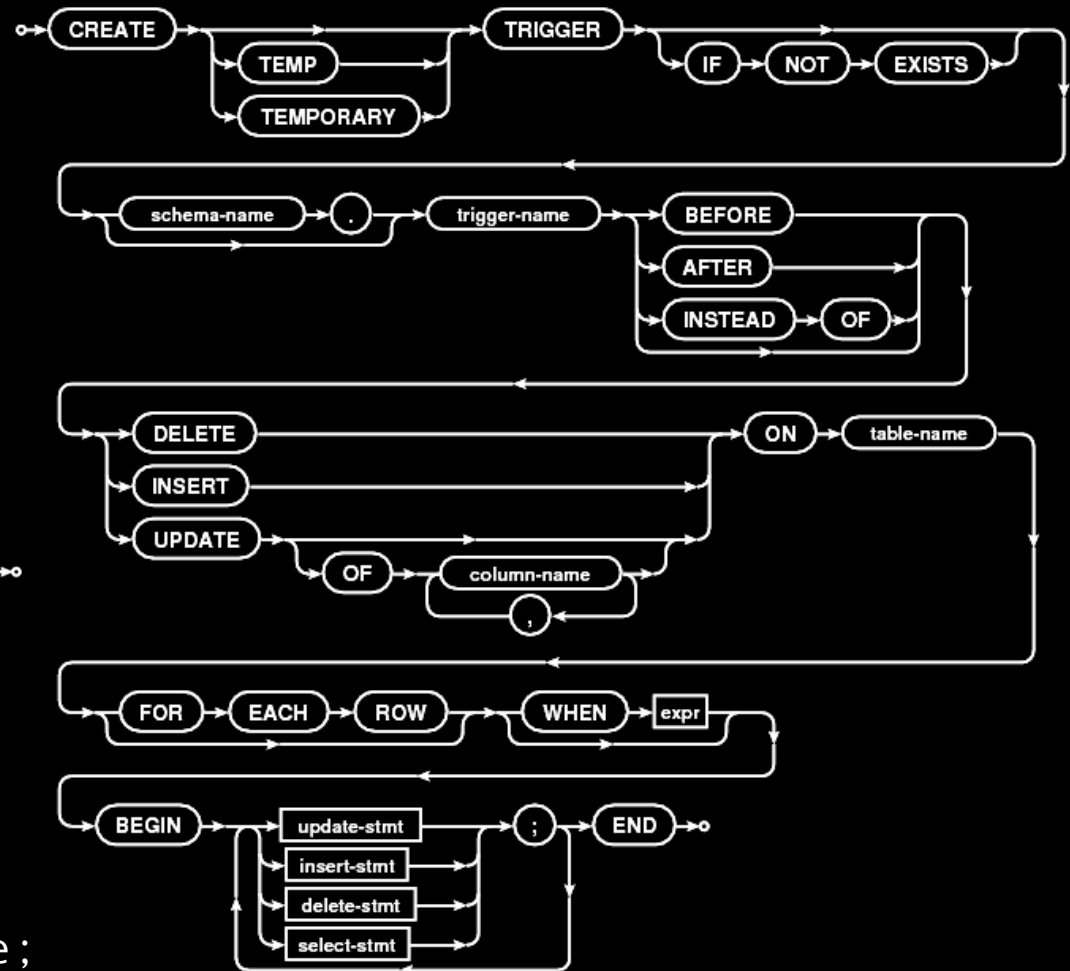
Conflict resolution algorithms

- **FAIL**: stop processing the rest of the current SQL statement
 - Do not undo any prior changes
 - If it fails on the 100th row, the actions taken due to the previous 99 rows are preserved
 - The transaction remains active (if within one)
- **ABORT**: stop processing the rest of the current SQL statement and abort
 - **Undo** any prior changes made by the **current SQL statement**
 - Changes caused by prior SQL statements within the same transaction are preserved
 - The transaction remains active
- **ROLLBACK**: stop processing the rest of the current SQL statement and rollback
 - **Undo** any prior changes made by **all SQL statements** in the transaction
 - **End the transaction**
 - If not within a transaction, ROLLBACK and ABORT are the same
- **IGNORE**: skip the one row that violates the constraint
 - Continue processing subsequent rows as if nothing went wrong
 - Do not return an error to the application

Triggers syntax



- Drop with:
DROP TRIGGER trigger_name;



Triggers in Postgres

- PostgreSQL supports triggers similarly to SQLite

```
CREATE [CONSTRAINT] TRIGGER name
  {BEFORE | AFTER | INSTEAD OF} {event [OR ...]}
  ON table_name
  [FROM referenced_table_name]
  [NOT DEFERRABLE | [DEFERRABLE] [INITIALLY IMMEDIATE | INITIALLY DEFERRED]]
  [REFERENCING { {OLD | NEW} TABLE [AS] transition_relation_name } [ ... ]
]
  [FOR [EACH] { ROW | STATEMENT } ]
  [WHEN (condition)]
  EXECUTE PROCEDURE function_name (arguments)
```

where *event* can be one of:

```
INSERT
UPDATE [ OF column_name [, ... ] ]
DELETE
TRUNCATE
```

```
CREATE FUNCTION function_name() RETURNS trigger AS $function_name$
BEGIN ... END
```

Transition Tables

- new/old iterate over rows only
- Can reference new/old rows of a *statement trigger* through **transition tables**

```
CREATE TRIGGER my_trigger
  BEFORE INSERT ON orig_table
  REFERENCING OLD TABLE AS old_tbl NEW TABLE AS new_tbl
  FOR EACH ROW
  EXECUTE PROCEDURE func();
```

```
CREATE FUNCTION func() RETURNS trigger AS $func$
BEGIN ... END
```

Constraint Triggers in Postgres

```
CREATE CONSTRAINT TRIGGER my_trigger  
...
```

- Must be AFTER ROW triggers
- The timing of the trigger firing can be adjusted
 - at the end of the statement causing the triggering event, or
 - at the end of the containing transaction (deferred)
- SET CONSTRAINTS
 - set constraint check timing for the current transaction
 - SET CONSTRAINTS { ALL | name [, ...] } { DEFERRED | IMMEDIATE }

(Non) Deferrable Constraint Triggers in Postgres

- **NOT DEFERRABLE**
 - Will be checked immediately after every command
- **DEFERRABLE**
 - Checking can be postponed until the end of the transaction
- **INITIALLY IMMEDIATE**
- **INITIALLY DEFERRED**
 - For deferrable constraints
 - Specify the default time to check the constraint
 - **INITIALLY IMMEDIATE**: check after each statement. This is the default.
 - **INITIALLY DEFERRED**: checked only at the end of the transaction
 - The constraint check time can be altered with the `SET CONSTRAINTS` command

Part II

Introduction to Transactions

Transaction Concept

- A **transaction** is a unit of program execution that accesses and possibly updates various data items
- Example: transaction to transfer \$50 from account A to account B:
 1. read(A)
 2. $A := A - 50$
 3. write(A)
 4. read(B)
 5. $B := B + 50$
 6. write(B)
- Two main issues to deal with:
 - Failures such as hardware failures, system crashes, query failures (e.g., triggers, conflicts)
 - Concurrent execution of multiple transactions

Atomicity requirement

- Transaction to transfer \$50 from account A to account B:
 1. read(A)
 2. $A := A - 50$
 3. write(A)
 4. read(B)
 5. $B := B + 50$
 6. write(B)
- What if the transaction fails at step 5 ?
 - Money will be “lost” leading to an inconsistent database state
 - Failure could be due to software or hardware
- The system should ensure that updates of a partially executed transaction are not reflected in the database

Durability requirement

- Transaction to transfer \$50 from account A to account B:
 1. `read(A)`
 2. `A := A - 50`
 3. `write(A)`
 4. `read(B)`
 5. `B := B + 50`
 6. `write(B)`
- The updates to the database by the transaction must persist even if there are software or hardware failures
- Once the user has been notified that the transaction has completed (i.e., the transfer of the \$50 has taken place) the state of the database should always reflect that

Consistency requirement

- Transaction to transfer \$50 from account A to account B:
 1. read(A)
 2. $A := A - 50$
 3. write(A)
 4. read(B)
 5. $B := B + 50$
 6. write(B)
- In above example: the sum of A and B is unchanged
- In general, consistency requirements include
 - Explicit integrity constraints, e.g., primary keys, foreign keys, unique values
 - Implicit integrity constraints, e.g., balances minus loans must equal cash-in-hand
- A transaction must see a consistent database
 - During transaction execution the database may be temporarily inconsistent
 - When the transaction completes successfully the database must be consistent

Isolation requirement

- Transaction to transfer \$50 from account A to account B:

User 1

```
1. read(A)
2. A := A - 50
3. write(A)
```

User 2

```
read(A), read(B), print(A+B)
```

```
4. read(B)
5. B := B + 50
6. write(B)
```

- User 2 should not be allowed to see the temporarily inconsistent database
 - The sum $A+B$ should not be incorrect, otherwise money appear to be “lost”
- Provide the illusion that transactions execute **serially**, i.e., one after the other
 - User 1 fully executes his transaction, then User 2 fully executes his transaction
 - ...or the other way around

ACID properties

- **Atomicity.** Either all operations of the transaction are properly reflected in the database, or none are
- **Consistency.** The execution of a transaction in isolation preserves the consistency of the database
- **Isolation.** Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions
 - Intermediate results must be hidden from the outside world
 - For every pair of transactions T_i and T_j , it appears to T_i that
 - Either T_i finished execution before T_j started, or
 - T_j finished execution before T_i started
- **Durability.** After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures

Transaction example: atomic updates - commit

```
SELECT OrderTotal FROM Orders WHERE OrderNumber IN (100, 101);  
3835.68  
1380.64
```

Start a transaction

BEGIN TRANSACTION;

UPDATE Orders SET OrderTotal=1.99 WHERE OrderNumber=100;

UPDATE Orders SET OrderTotal=1.99 WHERE OrderNumber=101;

```
SELECT OrderTotal FROM Orders WHERE OrderNumber IN (100, 101);  
1.99  
1.99
```

... check things; satisfied all is in order ...

COMMIT TRANSACTION;

make transaction updates persistent & end trans.

```
SELECT OrderTotal FROM Orders WHERE OrderNumber IN (100, 101);  
1.99  
1.99
```

Transaction example: atomic updates - rollback

```
SELECT OrderTotal FROM Orders WHERE OrderNumber IN (100, 101);  
3835.68  
1380.64
```

Start a transaction

```
BEGIN TRANSACTION;  
UPDATE Orders SET OrderTotal=1.99 WHERE OrderNumber=100;  
UPDATE Orders SET OrderTotal=1.99 WHERE OrderNumber=101;  
SELECT OrderTotal FROM Orders WHERE OrderNumber IN (100, 101);  
1.99  
1.99
```

... check things; realized you made a mistake ...

```
ROLLBACK TRANSACTION;
```

undo transaction updates & and transaction

```
SELECT OrderTotal FROM Orders WHERE OrderNumber IN (100, 101);  
3835.68  
1380.64
```


Transaction state

- **Active:** the initial state; the transaction stays in this state while it is executing
- **Failed:** after the discovery that normal execution can no longer proceed
- **Aborted:** after the transaction has been rolled back and the database is restored to its state prior to the start of the transaction.
 - Two options after it has been aborted:
 - Restart the transaction (can be done only if no internal logical error)
 - Kill the transaction
- **Partially committed:** after the final statement has been executed
- **Committed:** after successful completion

Named transactions

- **BEGIN . . . END** is one way to denote a transaction
 - END and COMMIT are the same: complete and exit the transaction
 - ROLLBACK: undo all changes and cancel the transaction
 - Subsequent SQL statements are not part of the transaction
 - BEGIN cannot be used within a transaction (i.e., no nesting)
- **SAVEPOINT** starts a transaction that is **named** and can be **nested**

Savepoints are similar to snapshots

- **SAVEPOINT TransactionName**
 - Create a new “mark” named **TransactionName** in the transaction timeline
 - “checkpoint” the database, i.e., takes a logical snapshot of it
- **ROLLBACK TO TransactionName**
 - Rewind the timeline back to a point just after the **TransactionName** mark
 - “restore the checkpoint”
- **RELEASE TransactionName**
 - Erase marks from the timeline without actually making any changes to the database
 - ...true for nested savepoints only
- **COMMIT**
 - Commits all outstanding transactions and leaves transaction stack empty
- **ROLLBACK**
 - Undo all changes and cancel all outstanding transactions

Savepoints

- **SAVEPOINT TransactionName**
 - Starts a transaction that is **named** and can be **nested**
- **ROLLBACK TO TransactionName**
 - Undo all changes until the beginning of TransactionName, and
 - Cancel all intervening savepoints, and
 - Restart the transaction with the name TransactionName
- **RELEASE TransactionName**
 - Remove all savepoints back to and including TransactionName
 - Cannot rollback to these savepoints anymore
 - If TransactionName is inner savepoint: no write back of modifications; COMMIT does that
 - If TransactionName is outermost savepoint, so that the transaction stack becomes empty, then RELEASE is the same as COMMIT
- **COMMIT**: commits all outstanding transactions and leaves transaction stack empty
- **ROLLBACK**: undo all changes and cancel all outstanding transactions

Example: atomically delete rows

SalesOrders.sqlite

```
select * from order_details;
```

	OrderNumber	ProductNumber	QuotedPrice	QuantityOrdered
1	1	1	1200	2
2	1	6	635	3
3	1	11	1650	4
4	1	16	28	1
5	1	21	55	3
6	1	26	121.25	5
7	1	40	174.6	6
8	2	27	24	4
9	2	40	180	4
10	3	1	1164	5

Commit transaction

Begin transaction

```
savepoint transaction1;  
delete from order_details where orderNumber=1;  
delete from order_details where orderNumber=2;  
select * from order_details;  
commit;
```

```
select * from order_details;
```

	OrderNumber	ProductNumber	QuotedPrice	QuantityOrdered
1	3	1	1164	5
2	3	6	615.95	5
3	3	11	1650	1
4				

Example: rollback attempt to atomically delete rows

SalesOrders.sqlite

```
select * from order_details;
```

	OrderNumber	ProductNumber	QuotedPrice	QuantityOrdered
1	3	1	1164	5
2	3	6	615.95	5
3	3	11	1650	1
4	3	16	28	2

Rollback and end transaction

Begin transaction

```
savepoint transaction2;  
delete from order_details where orderNumber=3;  
delete from order_details where orderNumber=4;  
select * from order_details;  
rollback;
```

```
select * from order_details;
```

	OrderNumber	ProductNumber	QuotedPrice	QuantityOrdered
1	3	1	1164	5
2	3	6	615.95	5
3	3	11	1650	1
4	3	16	28	2

Example: nested named transactions

SalesOrders.sqlite

```
select * from order_details;
```

	OrderNumber	ProductNumber	QuotedPrice	QuantityOrdered
1	3	1	1164	5
2	3	6	615.95	5
3	3	11	1650	1
4	3	16	28	2

```
T1:X savepoint transaction1;
```

```
T1:X delete from order_details where orderNumber=3;
```

```
T1:X savepoint transaction2;
```

```
T2:X delete from order_details where orderNumber=4;
```

```
T2:X select * from order_details;
```

```
T2:X rollback to transaction2;
```

```
T2:X select * from order_details;
```

```
T2:X rollback to transaction1;
```

```
T1:X rollback;
```

```
select * from order_details;
```

	OrderNumber	ProductNumber	QuotedPrice	QuantityOrdered
1	3	1	1164	5
2	3	6	615.95	5
3	3	11	1650	1
4	3	16	28	2

Transaction initiation

- Transactions can begin **explicitly**
 - [BEGIN | SAVEPOINT] . . . [END | COMMIT | ROLLBACK | RELEASE]
- Transactions can begin **implicitly**
 - Default on most databases: each SQL statement is wrapped in its own transaction
 - Transaction begins at SQL statement start
 - Transaction commits at the statement's final “;”
 - Starting up a server connection → transaction begin
 - pgAdmin 4: Configure → Query Tools → Options → Autocommit / Autorollback
 - Starting up DB Browser for SQLite → transaction begin
- Do **SELECT** statements require a transaction?
 - At the usual isolation level, a SELECT
 - Should not read uncommitted writes
 - Should not read writes from transactions that commit while the SELECT is running

Transactions in Postgres, MySQL

- PostgreSQL support transactions similarly to SQLite
- This includes support for **SAVEPOINT**