

MLDS 400 Lab 3

Monte Carlo Methods & MCMC

Huiyu Wu

10/9/2023



NORTHWESTERN
UNIVERSITY

From last time

- Baseball example
 - Aperiodic
- Interpretation of stationary distribution


Monte Carlo Methods

- Rely on random sampling to obtain numerical results
- Use randomness to solve problems that may be deterministic
- Common applications
 - Optimization
 - Numerical Integration
 - Generating draws from probability distributions

Monte Carlo Integration

- Goal: compute $I = \int_a^b h(x)dx$
- Monte Carlo Method
 - Let $h(x) = g(x)p(x)$, where $p(x)$ is a pdf defined on $[a, b]$
 - $I = \int_a^b h(x)dx = \int_a^b g(x)p(x)dx = E_p[g(x)]$
 - $x_1, \dots, x_n \sim p(x)$ iid
 - $I = E_p[g(x)] \approx \frac{1}{n} \sum_{i=1}^n g(x_i)$

Monte Carlo Integration Estimating Pi



- Circle inscribed in unit square
- Area of circle = $\frac{\pi}{4}$
- Area of square = 1
- $X = (X_1, X_2)$, drawn iid from $\text{unif}\left(-\frac{1}{2}, \frac{1}{2}\right)$
- $P(X_1^2 + X_2^2 \leq \frac{1}{4}) = \frac{\pi}{4}$
- $\pi \approx \frac{4}{n} \sum_{i=1}^n \mathbb{1}(x_{1i}^2 + x_{2i}^2 \leq \frac{1}{4})$

Monte Carlo Integration Estimating Pi in R

```
set.seed(400)  # Seed RNG  
n = 10000  
x1 = runif(n,-0.5,0.5)  
x2 = runif(n,-0.5,0.5)  
z = x1^2+x2^2  
Pi = 4*sum(z<=0.25)/n  
Pi
```

```
## [1] 3.1412
```

```
pi
```

```
## [1] 3.141593
```

Markov Inequality

- Statement: $P(Y \geq a) \leq \frac{E(Y)}{a}$
- Conditions: Y is non-negative and $a > 0$.
- What if: $Y = (X - \mu)^2$ and $a = \delta^{-1}\sigma^2$

Error Estimation

- Chebyshev Inequality: $P\left((X - \mu)^2 \geq \frac{\sigma^2}{\delta}\right) \leq \delta$
- This implies: $P\left((\text{error})^2 \geq \frac{\text{var}[p(x)]}{n\delta}\right) \leq \delta$
- Number of samples needed to n to meet tolerance with $(1 - \delta)$ confidence: $n \geq \frac{\text{var}[p(x)]}{(\text{tolerance})^2 \delta}$
- How many samples do we need to 99% confident that we can compute pi to 2 decimal places? Note: $\text{var}[p(x)] = \frac{1}{12^2}$

$$n \geq \frac{100 * 100^2}{144} \approx 6944$$

Markov Chain Monte Carlo (MCMC)


- What if $p(\cdot)$ is difficult to sample from?
- Idea: construct a Markov chain, where the equilibrium distribution is the desired distribution $p(\cdot)$.
- Detailed balance: $\pi_i p_{ij} = \pi_j p_{ji}$
- Bayes' Theorem: $p(x|y) = \frac{p(y|x)p(x)}{\int p(y|x)p(x)dx} \propto p(y|x)p(x)$

Rejection Method

Algorithm 1: Rejection Methods

- 1 Choose $q(\cdot)$, approximation of $p(\cdot)$
 - 2 Find M st $p(x) \leq Mq(x)$, $\forall x$
 - 3 Draw y from $q(\cdot)$
 - 4 Draw $u \sim \text{unif}(0, 1)$
 - 5 Compute acceptance ratio $\alpha = \frac{p(x)}{Mq(x)}$
 - 6 **if** $u \leq \alpha$ **then**
 - 7 Accept $x = y$
 - 8 **else**
 - 9 Reject: return to line 3, and try again
-

Intuition



Metropolis Algorithm

Algorithm 2: Metropolis Algorithm

```
1 Choose  $x_0$ 
2 Choose  $q(x|y)$  symmetric (i.e.,  $q(x|y) = q(y|x)$ )
3 Choose  $f(x) \propto p(x)$ 
4 for  $t = 0, 1, \dots$  do
5     Draw  $x'$  from  $q(\cdot|x_t)$ 
6     Compute acceptance ratio  $\alpha = \frac{f(x')}{f(x_t)}$ 
7     Draw  $u \sim \text{unif}(0, 1)$ 
8     if  $u \leq \alpha$  then
9         Accept  $x_{t+1} = x'$ 
10    else
11        Reject, set  $x_{t+1} = x_t$ 
```

Metropolis-Hastings Algorithm

Algorithm 3: Metropolis-Hastings Algorithm

```
1 Choose  $x_0$ 
2 Choose  $q(x|y)$  (general)
3 Choose  $f(x) \propto p(x)$ 
4 for  $t = 0, 1, \dots$  do
5   Draw  $x'$  from  $q(\cdot|x_t)$ 
6   Compute acceptance ratio  $\alpha = \frac{f(x')q(x_t|x')}{f(x_t)q(x'|x_t)}$ 
7   Draw  $u \sim \text{unif}(0, 1)$ 
8   if  $u \leq \alpha$  then
9     Accept  $x_{t+1} = x'$ 
10  else
11    Reject, set  $x_{t+1} = x_t$ 
```

MCMC

- Metropolis & Metropolis-Hastings algorithms generate sample points $\{x_0, \dots, x_t, \dots\}$
- These form a Markov chain as the probability α of transitioning from x_{t-1} to x_t does not depend on $\{x_0, \dots, x_{t-2}\}$
- Stationary distribution of Markov chain equals target distribution $p(\cdot)$

MCMC

- Typical implementation
 - Burn-in period k – samples for algorithm to approach desired distribution
 - Keep every m samples – thinning method to reduce auto-correlation
- Diagnostics
 - Plot of sample points: poor mixing and/or patterns vs. well mixing
 - Serial autocorrelations, time lag
 - Formal tests: Geweke test, Raftery-Lewis test

MCMC

- Choice of $q(\cdot|\cdot)$
 - Random-walk chain: new value $x' = x_t + z$,

$$q(x'|x_t) = r(x' - x_t) = r(z)$$

If $r(z) = r(-z)$, then $q(\cdot|\cdot)$ is symmetric.(ex: normal)

- Independent chain: new value x' drawn independent from previous point

$$q(x'|x_t) = s(x') \text{ for any distribution } s(x')$$

Generally not symmetric

Gibbs Sampling

- Simpler to sample from conditional than from joint.

Algorithm 4: Gibbs Sampler

```
1 Choose  $\mathbf{x}^{(0)} = (x_1^{(0)}, \dots, x_n^{(0)})$ 
2 for  $t = 0, 1, \dots$  do
3   for  $j = 0, \dots, n$  do
4     Draw  $x_j^{(t+1)}$  from  $p(\cdot | x_1^{(t+1)}, \dots, x_{j-1}^{(t+1)}, x_{j+1}^{(t)}, \dots, x_n^{(t)})$ 
```

MCMC Example

- Poisson regression
- No closed-form solution for optimal weights β^* , must be found using numerical methods

$$y_i | x_i \sim \text{Poisson} \left(\mu_i = \exp \left(x_i^T \beta \right) \right)$$
$$p(\beta | y, X) \propto \exp \left(\sum_{i=1}^n \left[-\exp \left(x_i^T \beta \right) + y_i x_i^T \beta \right] \right) \pi(\beta)$$

MCMC Example in R

- Puffin Dataset (38 observations at Great Island, Newfoundland)
 - Nest nesting frequency (burrows per 9 square meters)
 - Grass grass cover (percentage)
 - Soil mean soil depth (in centimeters)
 - Angle angle of slope (in degrees)
 - Distance distance from cliff edge (in meters)

```
library(LearnBayes)
```

```
puffin[1:5,]
```

| ## | Nest | Grass | Soil | Angle | Distance |
|------|------|-------|------|-------|----------|
| ## 1 | 16 | 45 | 39.2 | 38 | 3 |
| ## 2 | 15 | 65 | 47.0 | 36 | 12 |
| ## 3 | 10 | 40 | 24.3 | 14 | 18 |
| ## 4 | 7 | 20 | 30.0 | 16 | 21 |
| ## 5 | 11 | 40 | 47.6 | 6 | 27 |

MCMC Example in R

#Frequentist view

```
pfit = glm(Nest~Grass+Soil+Angle+Distance,poisson,data=puffin)
summary(pfit)
```

```
##
## Call:
## glm(formula = Nest ~ Grass + Soil + Angle + Distance, family = poisson,
##      data = puffin)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.3262  -1.2984  -0.6617   0.8119   2.5304
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  3.069973   0.452568   6.783 1.17e-11 ***
## Grass        0.005441   0.003104   1.753  0.07960 .
## Soil         0.033441   0.010822   3.090  0.00200 **
## Angle       -0.030077   0.010724  -2.805  0.00504 **
## Distance    -0.089399   0.010680  -8.371 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
##      Null deviance: 310.427  on 37  degrees of freedom
## Residual deviance:  68.765  on 33  degrees of freedom
## AIC: 183.38
##
## Number of Fisher Scoring iterations: 6
```

MCMC Example in R

```
#Bayesian method
```

```
library(MCMCpack)
```

```
#Assumes normal prior
```

```
Bpfit = MCMCpoisson(Nest+Grass+Soil+Angle+Distance,data=puffin,burnin=1000,mcmc=25000,thin=25)
```

```
summary(Bpfit)
```

```
##
```

```
## Iterations = 1001:25976
```

```
## Thinning interval = 25
```

```
## Number of chains = 1
```

```
## Sample size per chain = 1000
```

```
##
```

```
## 1. Empirical mean and standard deviation for each variable,
```

```
##    plus standard error of the mean:
```

```
##
```

```
##
```

```
##           Mean          SD Naive SE Time-series SE
```

```
## (Intercept) 3.066791 0.468017 1.480e-02      0.0164472
```

```
## Grass       0.005422 0.003108 9.827e-05      0.0001072
```

```
## Soil        0.033598 0.011199 3.541e-04      0.0004051
```

```
## Angle       -0.030116 0.011293 3.571e-04      0.0004012
```

```
## Distance    -0.090009 0.011227 3.550e-04      0.0004010
```

```
##
```

```
## 2. Quantiles for each variable:
```

```
##
```

```
##           2.5%       25%       50%       75%       97.5%
```

```
## (Intercept) 2.196294 2.734165 3.059824 3.364891 3.993143
```

```
## Grass       -0.000678 0.003212 0.005457 0.007519 0.011188
```

```
## Soil        0.010709 0.026503 0.034166 0.041186 0.055244
```

```
## Angle       -0.051478 -0.037771 -0.029952 -0.022732 -0.007356
```

```
## Distance    -0.111761 -0.097584 -0.089772 -0.082416 -0.067911
```