

# MLDS 400 Lab 8

## Data Imputation

Huiyu Wu

11/20/2023



NORTHWESTERN  
UNIVERSITY

# Missing Data

- Is data missing at random?
- Should we disregard missing or corrupt data? Or can we still use it?
- Are outliers truly outliers or bad data?

# Dataset

- Adjusted closing prices for Dow Jones members from 2018

```
library(quantmod); library(dplyr)
tickers = c("AAPL", "AMGN", "AXP", "BA", "CAT", "CRM", "CSCO",
            "CVX", "DIS", "KO", "GS", "HD", "HON", "IBM", "INTC",
            "JNJ", "JPM", "MCD", "MMM", "MRK", "MSFT", "NKE",
            "PG", "TRV", "UNH", "VZ", "V", "WMT", "WBA")
stocks = lapply(tickers, getSymbols, from="2019-01-01",
               to="2019-12-31", auto.assign=FALSE)
stocks = data.frame(Reduce(function(df1, df2)
  merge(df1, df2, by=0, all.x=T), stocks))
stocks = select(stocks, contains("Adjusted"))
colnames(stocks) = tickers
```

# Dataset

```
stocks[1:10,1:5]
```

##		AAPL	AMGN	AXP	BA	CAT
##	2019-01-02	37.89333	165.9931	89.08566	314.6451	111.8869
##	2019-01-03	34.11888	163.4673	87.34676	302.1006	107.5754
##	2019-01-04	35.57538	169.0552	91.28265	317.8226	113.4539
##	2019-01-07	35.49620	171.3301	91.77814	318.8234	113.5247
##	2019-01-08	36.17287	173.5272	92.22688	330.8919	114.8881
##	2019-01-09	36.78714	173.3196	92.39515	334.0985	115.3308
##	2019-01-10	36.90472	175.3178	91.86228	342.6300	117.7123
##	2019-01-11	36.54239	173.4840	92.13339	342.9118	116.9509
##	2019-01-14	35.99291	170.2143	92.03056	340.4437	116.6853
##	2019-01-15	36.72957	172.4633	91.60985	342.2705	115.7026

# Missing Data

```
library(missForest)
```

```
set.seed(400)
```

```
# randomly replace 10% of values with NAs
```

```
stocks.na = prodNA(stocks, noNA=0.1)
```

```
stocks.na[1:5,1:5]
```

##		AAPL	AMGN	AXP	BA	CAT
##	2019-01-02	37.89333	165.9931	NA	314.6451	111.8869
##	2019-01-03	34.11888	163.4673	87.34676	302.1006	107.5754
##	2019-01-04	35.57538	169.0552	NA	317.8226	NA
##	2019-01-07	35.49620	171.3301	91.77814	318.8234	113.5247
##	2019-01-08	36.17287	173.5272	92.22688	330.8919	114.8881

# Missing Data

```
# check missing data  
is.na(stocks.na[1:5,1:5])
```

```
##           AAPL  AMGN   AXP    BA    CAT  
## 2019-01-02 FALSE FALSE  TRUE  FALSE FALSE  
## 2019-01-03 FALSE FALSE FALSE  FALSE FALSE  
## 2019-01-04 FALSE FALSE  TRUE  FALSE  TRUE  
## 2019-01-07 FALSE FALSE FALSE  FALSE FALSE  
## 2019-01-08 FALSE FALSE FALSE  FALSE FALSE
```

# Handling Missing Data

```
AAPL.na = stocks.na$AAPL  
AAPL.na[1:10]
```

```
## [1] 37.89333 34.11888 35.57538 35.49620 36.17287 36.78714 36.90472 36.54239  
## [9]      NA 36.72957  
# disregard missing data  
mean(AAPL.na, na.rm=T)
```

```
## [1] 50.31591  
# remove missing observations  
AAPL.omit = na.omit(AAPL.na)  
AAPL.omit[1:10]
```

```
## [1] 37.89333 34.11888 35.57538 35.49620 36.17287 36.78714 36.90472 36.54239  
## [9] 36.72957 37.17828
```

# Random Sampling

```
random.imp = function (a){  
  missing = is.na(a)  
  n.missing = sum(missing) # number of missing values  
  a.obs = a[!missing]  
  imputed = a  
  # sample with replacement  
  imputed[missing] = sample(a.obs, n.missing, replace=T)  
  return(imputed)  
}  
AAPL.rndimp = random.imp(AAPL.na)  
AAPL.rndimp[1:10]
```

```
## [1] 37.89333 34.11888 35.57538 35.49620 36.17287 36.78714 36.90472 36.54239  
## [9] 43.45668 36.72957
```



# Most Common Value

```
x = c(1,1,NA,3,4,4,5,5,5,5,6,NA)
# compute the mode
Mode = function(x) {
  mode = as.numeric(names(sort(table(x),decreasing=T))[1])
  return(mode)
}
Mode(x)
```

```
## [1] 5
```

## Most Common Value

```
mcv.imp = function (a){  
  missing = is.na(a)  
  imputed = a  
  imputed[missing] = Mode(a)  
  return(imputed)  
}  
x.mcv = mcv.imp(x)  
x.mcv
```

```
## [1] 1 1 5 3 4 4 5 5 5 5 6 5
```

# Average Value

```
avg.imp = function (a){  
  missing = is.na(a)  
  imputed = a  
  imputed[missing] = mean(a, na.rm=T)  
  return(imputed)  
}  
AAPL.avgimp = avg.imp(AAPL.na)  
AAPL.avgimp[1:10]
```

```
## [1] 37.89333 34.11888 35.57538 35.49620 36.17287 36.78714 36.90472 36.54239  
## [9] 50.31591 36.72957
```

# Last Value

```
AAPL.last = na.locf(AAPL.na)  
AAPL.last[1:10]
```


```
## [1] 37.89333 34.11888 35.57538 35.49620 36.17287 36.78714 36.90472 36.54239  
## [9] 36.54239 36.72957
```

# Linear


```
n = length(AAPL.na)
AAPL.linear = approxfun(1:n, AAPL.na)(1:n)
AAPL.linear[1:10]
```

```
## [1] 37.89333 34.11888 35.57538 35.49620 36.17287 36.78714 36.90472 36.54239
## [9] 36.63598 36.72957
```

## k-Nearest Neighbors (k-NN)



# k-NN



# Imputation with 1-NN

```
library(DMwR)
```

```
stocks.1NN = knnImputation(stocks.na,k=1)  
stocks.1NN[1:5,1:5]
```

##		AAPL	AMGN	AXP	BA	CAT
##	2019-01-02	37.89333	165.9931	91.86228	314.6451	111.8869
##	2019-01-03	34.11888	163.4673	87.34676	302.1006	107.5754
##	2019-01-04	35.57538	169.0552	91.86228	317.8226	117.7123
##	2019-01-07	35.49620	171.3301	91.77814	318.8234	113.5247
##	2019-01-08	36.17287	173.5272	92.22688	330.8919	114.8881



# Imputation with 5-NN

```
stocks.5NN = knnImputation(stocks.na,k=5)  
stocks.5NN[1:5,1:5]
```

##		AAPL	AMGN	AXP	BA	CAT
##	2019-01-02	37.89333	165.9931	92.17218	314.6451	111.8869
##	2019-01-03	34.11888	163.4673	87.34676	302.1006	107.5754
##	2019-01-04	35.57538	169.0552	91.95454	317.8226	117.7144
##	2019-01-07	35.49620	171.3301	91.77814	318.8234	113.5247
##	2019-01-08	36.17287	173.5272	92.22688	330.8919	114.8881

# Multivariate Imputation via Chained Equations

- Multivariate Imputation via Chained Equations (MICE)
- “Generates multiple imputations for incomplete multivariate data by Gibbs sampling.”
  - Goes column-by-column, imputing missing values, then repeats.
- Default methods:
  - Predictive mean matching (PMM) - numeric data
  - Logistic regression - binary factor
  - Polytomous logistic regression - unordered factor ( $>2$  levels)
  - Proportional odds model - ordered factor ( $>2$  levels)
- Documentation:  
<https://cran.r-project.org/web/packages/mice/mice.pdf>

# PMM


- 1 Estimate the linear regression model using  $y, x$ , obtaining parameters  $\{\hat{\beta}, \Sigma_{\beta}\}$ .
- 2 Draw  $\beta^* \sim \mathcal{N}(\hat{\beta}, \Sigma_{\beta})$ .
- 3 Using  $\beta^*$ , generate predicted  $\hat{y}$  for all samples (both missing and not).
- 4 For each missing value  $y$ , identify the set of samples with observed  $y$  whose predicted value  $\hat{y}$  is close to the predicted value of the missing sample.
- 5 Randomly pick observed value  $y^*$  from the set from Step 4 and assign it as the imputed missing value.
- 6 Repeat.

# MICE

```
library(mice)
```

```
md.pattern(stocks.na[1:5,1:5])
```

AAPL AMGN BA CAT AXP



```
## AAPL AMGN BA CAT AXP
## 3 1 1 1 1 1 0
## 1 1 1 1 1 0 1
## 1 1 1 1 0 0 2
## 0 0 0 1 2 3
```

# MICE

```
stocks.mice = mice(stocks.na, seed=400, print=F)
stocks.mice$imp$AAPL[1:5]
```

##		1	2	3	4	5
##	2019-01-14	35.57538	35.57538	37.89333	36.54239	35.49620
##	2019-01-28	37.39903	36.64078	37.62939	36.64078	36.93352
##	2019-02-13	41.18138	41.66816	41.22236	41.80937	41.45610
##	2019-02-15	41.18138	42.16459	41.45610	42.14049	42.01036
##	2019-04-01	45.99916	46.18543	46.18543	45.63871	46.97647
##	2019-04-03	48.34578	46.97647	46.18543	44.95034	48.33358
##	2019-04-16	50.14175	50.40135	49.19237	48.01081	49.12656
##	2019-04-22	48.75693	49.16334	50.40135	49.12656	48.75693
##	2019-05-13	45.48291	45.47809	41.98627	45.48291	43.78881
##	2019-06-06	42.90998	45.01299	43.78881	45.14038	43.59602
##	2019-06-20	49.92180	50.14418	50.23989	49.44879	48.95064
##	2019-07-02	50.14175	49.64232	49.16334	49.12656	48.95064
##	2019-07-12	49.46875	50.74397	49.30729	51.02790	48.80531
##	2019-07-30	52.62104	50.68812	50.75368	49.99891	51.78572

# MICE

```
# get the 5th imputation  
stocks.miceImp = complete(stocks.mice, 5)  
stocks.miceImp[1:5,1:5]
```

##		AAPL	AMGN	AXP	BA	CAT
##	2019-01-02	37.89333	165.9931	92.13339	314.6451	111.8869
##	2019-01-03	34.11888	163.4673	87.34676	302.1006	107.5754
##	2019-01-04	35.57538	169.0552	92.39515	317.8226	114.2664
##	2019-01-07	35.49620	171.3301	91.77814	318.8234	113.5247
##	2019-01-08	36.17287	173.5272	92.22688	330.8919	114.8881