

MLDS 422 – Fall 2023

Project 3

Due Wednesday, 11/29/23 at 11:59pm

Exploratory Data Analysis

Write Python code to do the following with the *movie_reviews.xlsx* spreadsheet:

1. Find the number of movies that were released in theatres each year.
 - There are some erroneous values in the data
 - Visualize the number of movies each year and describe what you see
 - Explain what assumptions and fixes you plan to make in the data as you move forward with your analysis
2. Looking only at movies that were released in theatres before 2010:
 - Find the 5 highest rated movies by critics (*critic_rating*)
 - Find the 5 highest rated movies by the general audience (*audience_rating*)
 - Take a look at the movie titles – these top movies don't seem very popular
3. Create a new DataFrame containing only “popular” movies that were released in theatres before 2010.
 - Assume popular movies are those with more audience reviews than the average number of audience reviews of all movies before 2010
 - Find the 5 highest rated movies by critics (*critic_rating*)
 - Find the 5 highest rated movies by the general audience (*audience_rating*)
 - Take a look at the movie titles – these top movies should make more sense
4. Using your new DataFrame (popular movies released in theatres before 2010), answer the following questions about the *rating* column:
 - What percent of movies fall under each type of rating (R, PG-13, etc.)? What are your takeaways from the summary table?
 - Create a visualization that shows for each rating (R, PG-13, etc.), the average critic rating and the average audience rating. What are your takeaways from the visualization?
5. Create a pair plot of the new DataFrame. What are 3 insights you can take away from looking at the pair plot?
6. Using either pandas or data visualizations, find 3 more insights using any columns in the new DataFrame.

Feature Engineering

Your goal is to create a predictive model that will predict the `critic_rating`. Using the full data set that you originally read into Python (with the erroneous years fixed):

7. Split the data into a training and test set, with the training data including movies released in theatres before 2010 and the test data including movies released in theatres in 2010 and after.
8. If your goal is to predict the `critic_rating` before the first critic or audience rating gets posted for a movie, which columns in the data should you NOT use to create features? Update your training and test data sets to NOT include these columns.
9. Using only the training data, create a new DataFrame containing the following ID column and features:
 - `movie_title`
 - `runtime_in_minutes`
 - NEW: `kid_friendly` (1 if G or PG, 0 if other ratings)
 - NEW: dummy variable columns for each genre
10. Create 3 new features that you think will do a good job predicting the `critic_rating`. Each new feature should use various combinations of the columns from your training data.

Modeling

11. Make sure you apply the same transformations on your `X_test` and `y_test` data sets that you applied on the `X_train` and `y_train` data sets.
12. Make sure that your `X_train`, `y_train`, `X_test` and `y_test` data sets only contain columns of numeric and non-null values. Explain and justify how you decide to deal with data issues.
13. Fit 3 linear regression models on the training data:
 - Model 1: Use only `runtime_in_minutes`
 - Model 2: Use `runtime_in_minutes` and `kid_friendly`
 - Model 3: Use `runtime_in_minutes`, `kid_friendly` and the dummy columns for the genres
14. Score the linear regression models on the test data by writing a function where you can input the `y_test` and `y_pred` values (`y_pred` = predicted values after you apply the fitted model to your `X_test` data), and it outputs the following metrics: R^2 , MAE and RMSE. Apply the function to the three models that you've fit so far.

15. Which model performs the best so far? Which features seem to do a good job predicting the critic rating (hint: you can check p-values using statsmodels)?
16. Try fitting 3 more linear regression models on your own using a combination of the columns so far (runtime_in_minutes, kid-friendly and the dummy columns for genre) and your newly engineered features.
- Each subsequent model should attempt to do a better job at prediction than the previous model (even if the metrics don't end up being better, your choices should make sense)
 - With each new model, explain why are you making your decisions of which feature(s) to include / remove
17. Out of the 6 models you created, which model performs the best? Which features seem to do a good job predicting the critic rating?
18. List 3 other things you could do at this point to try and improve your model.